

Traditional Programming



Machine Learning



Introduction to Machine Learning

Outline

- The main machine learning tasks
- Good generalization
- Naïve Bayes classifier
- k-Nearest Neighbour classifier
- Classification errors (training, validation and test)

Reference

Part 5 of AIMA textbook (4th Ed)
Chapter 19 “Learning from Examples”


Machine Learning Tasks

Examples of Statistical Learning Problems

- Predict whether someone will have a heart attack on the basis of demographic, diet and clinical measurements
- Customize an email spam detection system

	george	you	hp	free	!	edu	remove
spam	0.00	2.26	0.02	0.52	0.51	0.01	0.28
email	1.27	1.27	0.90	0.07	0.11	0.29	0.01





- Identify the numbers in a handwritten postcode



- Identify emergency landing places for UAV's

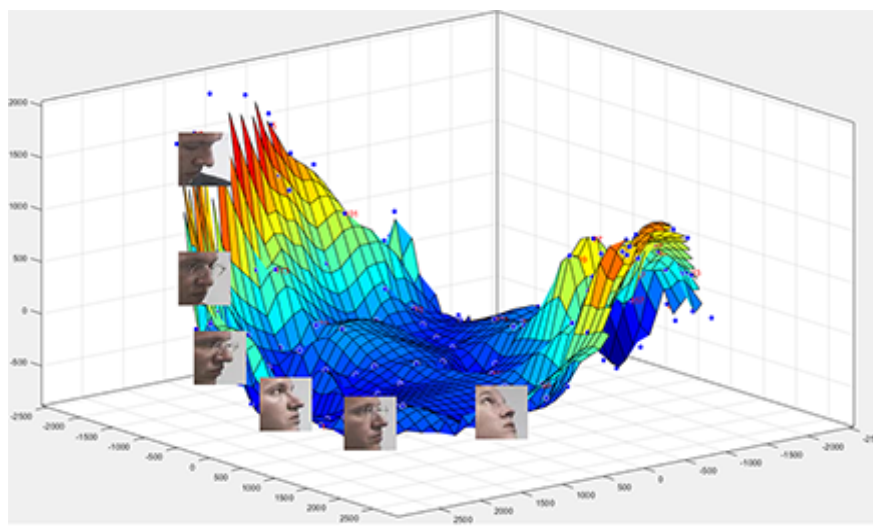



Machine Learning Tasks

- Classification (label prediction)
- Regression (function approximation)
- Clustering
- Probability density estimation
- Policy learning (learning from experience)

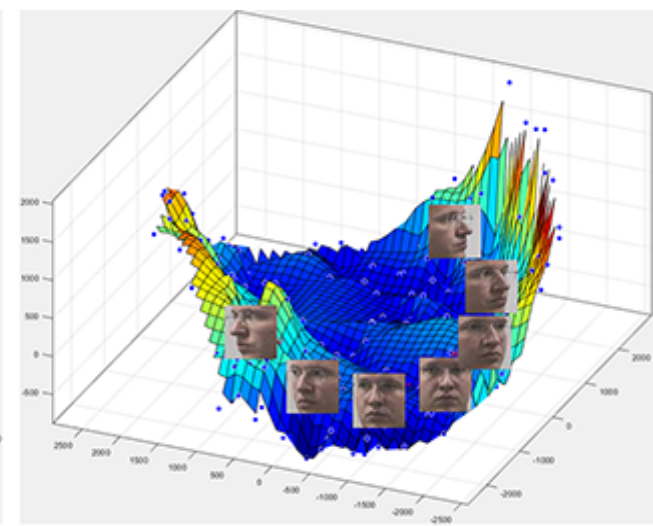


Manifold Learning

- Noise removal
- Classification
- Autoencoders
- Embeddings



(a)



(b)

Unsupervised Learning

- No outcome variable, just a set of predictors (features) measured on a set of samples
- Objective is more fuzzy — **find groups** of samples that behave similarly, find features that behave similarly, find linear combinations of features with the most variation
- Difficult to know how well you are doing!
- Different from supervised learning, but can be useful as a pre-processing step for supervised learning


Supervised Learning

- Outcome measurement Y (also called dependent variable, response, target)
- Vector of p predictor measurements X (also called inputs, regressors, covariates, features, independent variables)
- In the regression problem, Y is quantitative (e.g price, blood pressure)
- In the classification problem, Y takes values in a finite, unordered set (survived/died, digit 0-9, safe for landing)
- We have training data $(x_1, y_1), \dots, (x_N, y_N)$. These are observations (examples, instances) of these measurements

ML Objectives

- Accurately **predict the labels of new** test cases
- Understand which inputs affect the outcome, and how
- Assess the quality of predictions and inferences

Polynomial Curve Fitting




Let's look at **generalization**!
Good generalization is what we care about

Blue **dots** are our data points
They are noisy measurements of hidden/unknown **green** data

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$


Sum-of-Squares Error Function



$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

$E(\mathbf{w})$ measures the discrepancy between the model $y(x, \mathbf{w})$ and the dataset $\{(x_1, y_1), \dots, (x_N, y_N)\}$


0th Order Polynomial



M is the degree of the polynomial

In red, the best fit for a polynomial model of degree $M=0$


1st Order Polynomial



As M increases, the capacity of the model increases.
Capacity \sim ability of the model to fit the data


In red, the best fit for a polynomial of degree $M=1$

3rd Order Polynomial




As M increases, the capacity of the model increases

9th Order Polynomial



At some point the capacity can become too large!

Over-fitting



Only the examples in the **Training Set** are used for selecting w . The examples of the **Test Set** are used for the evaluation of the generalization performance of the trained model.

Root-Mean-Square (RMS) Error:

$$E_{\text{RMS}} = \sqrt{2E(\mathbf{w}^*)/N}$$

Polynomial Coefficients


	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Weight explosion as M increases

Data Set Size:

$$N = 15$$

9th Order Polynomial



The appropriate capacity of the model (M) depends on the size of the training set

Regularization

Idea : Penalize large coefficient values


$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

We can control the capacity of the parameterized function by imposing a penalty for large weights

Regularization:

vs. E_{RMS}

$\ln \lambda$




Small λ value
for the regularization
term in the loss

Large λ value
for the regularization
term in the loss

Probability Refresher

Probability Theory

Apples and Oranges



2-step random process:
(1) select randomly a bag
(2) select randomly a fruit from that bag

Concept review:

- **Random variable**
- **Conditional probability**

Random variable examples:

B : colour of the bag

F : type of fruit

Question:

$$P(B|F) ?$$

What is the probability that the fruit was picked from the blue bag given that it is an orange?

The Rules of Probability


Sum Rule

$$p(X) = \sum_Y p(X, Y)$$

Product Rule

$$p(X, Y) = p(Y|X)p(X)$$

Y \ X	x ₁	x ₂	x ₃	x ₄	p _y (Y) ↓
y ₁	4/32	2/32	1/32	1/32	8/32
y ₂	2/32	4/32	1/32	1/32	8/32
y ₃	2/32	2/32	2/32	2/32	8/32
y ₄	8/32	0	0	0	8/32
p _x (X) →	16/32	8/32	4/32	4/32	32/32



Bayes' Theorem

$$p(Y|X) = \frac{p(X|Y)p(Y)}{p(X)}$$

$$p(X) = \sum_Y p(X|Y)p(Y)$$


posterior likelihood \times prior

Y: class

X: measurements

Probability Densities

*For non discrete
random variables*



$$p(x \in (a, b)) = \int_a^b p(x) dx$$

$$P(z) = \int_{-\infty}^z p(x) dx$$

$$p(x) \geq 0$$

$$\int_{-\infty}^{\infty} p(x) dx = 1$$

Classification Learning

- Given a training sample, we want to predict $y \in Y$ for a new $x = (x_1, \dots, x_n) \in X$
- Error minimised by $y = \operatorname{argmax}_y P(y | x_1, \dots, x_n)$
- Can estimate probabilities using observed frequencies

$$P(W) \simeq F(W)$$
$$P(W|Z) \simeq \frac{F(W, Z)}{F(Z)}$$

Estimating Probabilities

To estimate the probability of an attribute-value $A = v$ for a given class C we use

the *relative frequency* n_c / n

where n_c is the number of training instances that belong to the class C and have value v for the attribute A , and n is the number of training instances of the class C

Naïve Bayes Classifier

Let each instance x of a training set D be described by a conjunction of n attribute values $\langle a_1, a_2, \dots, a_n \rangle$ and let the target function, be such that it can take any value from a finite set V .

$$\begin{aligned} v_{\text{MAP}} &= \operatorname{argmax}_{v_j \in V} P(v_j | a_1, \dots, a_n) \\ &= \operatorname{argmax}_{v_j \in V} \frac{P(a_1, \dots, a_n | v_j) P(v_j)}{P(a_1, \dots, a_n)} \\ &= \operatorname{argmax}_{v_j \in V} P(a_1, \dots, a_n | v_j) P(v_j) \\ &= \operatorname{argmax}_{v_j \in V} P(v_j) \prod_i P(a_i | v_j) \end{aligned}$$



Naïve Bayes simplifying assumption is that attributes are conditionally independent!

Example

Consider the weather data and we have to classify the instance:

< Outlook = sunny, Temp = cool, Hum = high, Wind = strong >

The task is to predict the value (*yes* or *no*) of the concept **PlayTennis**. We apply the Naïve Bayes rule:

$$\begin{aligned} v_{MAP} &= \operatorname{argmax}_{v_j \in \{yes, no\}} P(v_j) \prod_i P(a_i | v_j) \\ &= \operatorname{argmax}_{v_j \in \{yes, no\}} P(v_j) P(\text{Outlook} = \text{sunny} | v_j) P(\text{Temp} = \text{cool} | v_j) \\ &\quad P(\text{Hum} = \text{high} | v_j) P(\text{Wind} = \text{strong} | v_j) \end{aligned}$$

Example: Estimating Probabilities

Outlook	Temperature	Humidity	Windy	Class
sunny	hot	high	false	N
sunny	hot	high	true	N
overcast	hot	high	false	P
rain	mild	high	false	P
rain	cool	normal	false	P
rain	cool	normal	true	N
overcast	cool	normal	true	P
sunny	mild	high	false	N
sunny	cool	normal	false	P
rain	mild	normal	false	P
sunny	mild	normal	true	P
overcast	mild	high	true	P
overcast	hot	normal	false	P
rain	mild	high	true	N

$$P(\text{yes}) = 9/14$$

$$P(\text{no}) = 5/14$$

Outlook

$$P(\text{sunny}|\text{yes}) = 2/9$$

$$P(\text{sunny}|\text{no}) = 3/5$$

$$P(\text{overcast}|\text{yes}) = 4/9$$

$$P(\text{overcast}|\text{no}) = 0$$

$$P(\text{rain}|\text{yes}) = 3/9$$

$$P(\text{rain}|\text{no}) = 2/5$$

Temp

$$P(\text{hot}|\text{yes}) = 2/9$$

$$P(\text{hot}|\text{no}) = 2/5$$

$$P(\text{mild}|\text{yes}) = 4/9$$

$$P(\text{mild}|\text{no}) = 2/5$$

$$P(\text{cool}|\text{yes}) = 3/9$$

$$P(\text{cool}|\text{no}) = 1/5$$

Hum

$$P(\text{high}|\text{yes}) = 3/9$$

$$P(\text{high}|\text{no}) = 4/5$$

$$P(\text{normal}|\text{yes}) = 6/9$$

$$P(\text{normal}|\text{no}) = 1/5$$

Windy

$$P(\text{true}|\text{yes}) = 3/9$$

$$P(\text{true}|\text{no}) = 3/5$$

$$P(\text{false}|\text{yes}) = 6/9$$

$$P(\text{false}|\text{no}) = 2/5$$

Example

$$P(\text{yes}) P(\text{sunny}|\text{yes}) P(\text{cool}|\text{yes}) P(\text{high}|\text{yes}) P(\text{strong}|\text{yes}) = .0053$$

$$P(\text{no}) P(\text{sunny}|\text{no}) P(\text{cool}|\text{no}) P(\text{high}|\text{no}) P(\text{strong}|\text{no}) = .0206$$

Thus, the naïve Bayes classifier assigns the value *no* to
PlayTennis!



Nearest Neighbours Classifiers

Instance-Based Methods

- Instance-based learning:
 - Store training examples and delay the processing (“lazy evaluation”) until a new instance must be classified
- Typical approaches
 - *k*-nearest neighbour approach
 - Instances represented as points in a Euclidean space.
 - Locally weighted regression
 - Constructs local approximation

The k -Nearest Neighbour Algorithm

- All instances correspond to points in the n -D space.
- The nearest neighbor are defined in terms of Euclidean distance.
- The target function could be discrete- or real- valued.
- For discrete-valued, the k -NN returns the most common value among the k training examples nearest to xq .
- **Voronoi diagram**: the decision surface induced by 1-NN for a typical set of training examples.



Fuzzy flavour

Discussion on the k -NN Algorithm

- The k -NN algorithm for continuous-valued target functions
 - Calculate the mean values of the k nearest neighbours
- Distance-weighted nearest neighbor algorithm
 - Weight the contribution of each of the k neighbours according to their distance to the query point x_q
 - giving greater weight to closer neighbours
 - Similarly, for real-valued target functions
- Robust to noisy data by averaging k -nearest neighbours
- Curse of dimensionality: distance between neighbours could be dominated by irrelevant attributes.
 - To overcome it, axes stretch or elimination of the least relevant attributes.

$$w = \frac{1}{d(x_q, x_i)^2}$$

Training, Validation and Test Errors

What is the difference between a *training set*, a *validation set* and a *test set*?

- The **training set** is used to select the parameter vector of the classifier
- The **validation set** is used to select the **capacity/complexity** of the classifier. For example, the capacity can be the maximum depth of a decision tree. Shallow decision trees have less discriminative power than deeper ones. But deeper decision trees might overfit the data.
- Assume that the capacity of your classifier is an integer n . In order to select the capacity of your classifier, you can train (using the training set) a number of classifiers with different capacity n , and evaluate the generalization of the classifier using the validation set. You should pick the capacity n for which the generalization error is the lowest.
- So, why do we need the **test set**?!
- The validation set cannot be used as a test set, because it is "tainted". It has been used to select the capacity/architecture of the classifier. So, it cannot be said that it is "unseen data". This is why, statisticians require the use of a different set (the test set) to assess the generalization error of the trained classifier).