

Classifier construction in sklearn and keras

Last modified on Saturday, 7 May 2022 by f.maire@qut.edu.au

In this exercise sheet, you will build and train different classifiers for the classification of handwritten digits. You will also use a number of metrics to assess the performance of your classifiers.

Exercise 1 – Retrieve and inspect the MNIST dataset

Use the scaffolding code in the provided file *sklearn_nn.py*, load the MNIST dataset.

```
from sklearn.datasets import fetch_openml
mnist = fetch_openml('mnist_784', version=1)
```

- Print the keys of the *mnist* object (which is an instance of a dictionary-like class called “Bunch”).
- What are the entries “data” and “target”?
- What are their shapes?

For convenience, alias the inputs and the class labels with

```
X, y = mnist["data"], mnist["target"]
```

Plot the example indexed 99 with the following code

```
import matplotlib as mpl
import matplotlib.pyplot as plt
some_digit = X[99]
some_digit_image = some_digit.reshape(28, 28)
plt.imshow(some_digit_image, cmap = mpl.cm.binary, interpolation="nearest")
plt.axis("off")
plt.show()
```

Exercise 2 – Linear classifier

Let's keep the same X and y from the previous exercise.

- What is the type of the entries of y ?
- Convert the entries of y to *uint8* (hint: use ndarray method 'astype')
- Create a training and test set using
 $X_{train}, X_{test}, y_{train}, y_{test} = X[:60000], X[60000:], y[:60000], y[60000:]$

Let's build a binary classifier to distinguish between '5' and 'non 5'

- Create boolean arrays (masks) to identify the digits that correspond to a '5' in the arrays y_{train} and y_{test} .

We will build a classifier for the digit class '5' that implements a regularized linear model with stochastic gradient descent (SGD) learning.

- Create an instance sgd_clf of the class *sklearn.linear_model.SGDClassifier*
- Train sgd_clf to discriminate between '5' and 'non 5' (hint: use the *fit* method)
- Predict the classes of the batch $X_{test}[50:80]$ with sgd_clf . Check the results by plotting the digits.
- Convert the entries of y to *uint8* (hint: use ndarray method 'astype')

A good way to evaluate a model is to use **cross-validation**. Let's use the *cross_val_score* function to evaluate your *SGDClassifier* model using K-fold cross-validation, with three folds. Remember that K-fold cross-validation means splitting the training set into K-folds (in this case, three), then making predictions and evaluating them on each fold using a model trained on the remaining folds.

- Import the function *cross_val_score* from the module *sklearn.model_selection*
- Call the *cross_val_score* with the appropriate parameters (you can use `scoring="accuracy"`)

Exercise 3 – Confusion Matrix

A good way to assess the performance of a classifier is to look at the **confusion matrix**. The general idea is to count the number of times instances of class A are classified as class B. For example, to know the number of times the classifier confused images of 5s with 3s, you would look in the 5th row and 3rd column of the confusion matrix.

To compute the confusion matrix, you first need to have a set of predictions, so they can be compared to the actual targets. You could make predictions on the test set, but let's keep it untouched for now (remember that you want to use the test set only at the very end of your project, once you have a classifier that you are ready to launch). Instead, you can use the `cross_val_predict` function from the module `sklearn.model_selection`.

Just like the `cross_val_score` function, `cross_val_predict` performs K-fold cross-validation, but instead of returning the evaluation scores, it returns the predictions made on each test fold. This means that you get a clean prediction for each instance in the training set ("clean" meaning that the prediction is made by a model that never saw the data during training).

- Compute the confusion matrix by combining the function `cross_val_predict` and `confusion_matrix`.
- What do the rows and columns of the confusion matrix represent?

Exercise 4 – Precision and Recall

The confusion matrix gives you a lot of information, but sometimes you may prefer a more concise metric. An interesting one to look at is the accuracy of the positive predictions; this is called the **precision** of the classifier.

$$\text{Precision} = TP / (TP + FP)$$


where TP is the number of true positives, and FP is the number of false positives.

Precision is typically used along with another metric named **recall**, also called sensitivity or true positive rate (TPR): this is the ratio of positive instances that are correctly detected by the classifier

$$\text{Recall} = TP / (TP + FN)$$

where FN is of course the number of false negatives.

- Compute the precision and recall of your classifier using appropriate functions from the `sklearn.metrics` module



Exercise 5 – Random Forest classifiers

Decision trees are still a popular type of classifiers when the inputs have discrete values. This week you can use them as a black box.

A **random forest** is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting. Random Forest classifiers can directly classify instances into multiple classes. You can call `predict_proba()` to get the list of probabilities that the classifier assigned to each instance for each class. Let's train a `RandomForestClassifier`. The `predict_proba()` method returns an array containing a row per instance and a column per class, each containing the probability that the given instance belongs to the given class.

- Create a random forest classifier by instantiating an object *forest_clf* of the class `RandomForestClassifier` (hint: look at the *sklearn.ensemble* module).
- Train the forest on the 10 digit classes (multiclass classification).
- Print the confusion matrix

Exercise 6 – Multi-layer Neural Network in sklearn

The `MLPClassifier` and `MLPRegressor` are sklearn implementations of **neural networks**. In this exercise, you will be training a neural network based multiclass classifier for the MNIST dataset (downloaded in Exercise 1). As neural networks are sensitive to the scale of the input, we will preprocess the data.

- Scale the data with a *StandardScaler* from the *sklearn.preprocessing* module
- Build a classifier *mlp* using the class *MLPClassifier* from the *sklearn.neural_network* module. Suggestion use “`hidden_layer_sizes=[10,10,10]`, `verbose=True`, `max_iter=1000`” as the parameter list of the constructor.
- Fit *mlp* to the train data and plot the confusion matrix (don't forget to apply the `StandardScaler` to *X_test*!)

Exercise 7 – Multi-layer Neural Network in Keras

In this exercise, use Keras to implement and train the neural network described in Exercise 6. Some scaffolding code is provided in the file *keras_nn.py*