

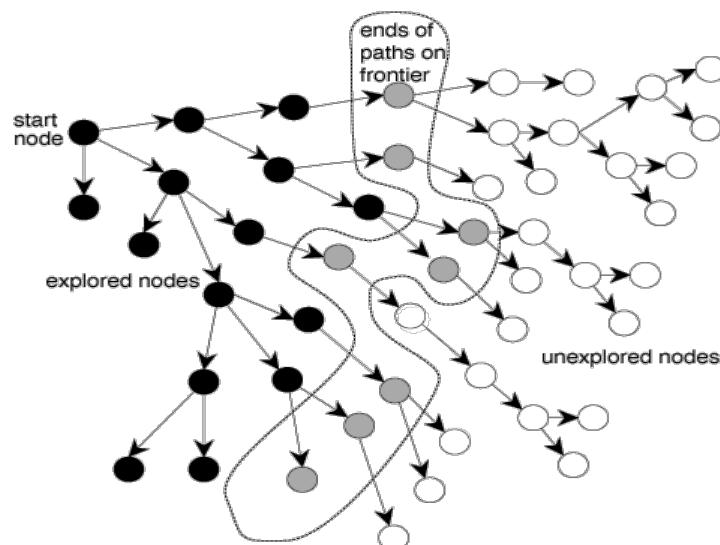
A* Tree Search

Informed Search



Credits: most material borrowed from AIMA or Dan Klein and Pieter Abbeel

Last modified on 2020/08/01 by f.maire@qut.edu.au



Reading for this week

- Chapter 3, Sections 3.5 to 3.6 of AIMA
 - Russell and Norvig Textbook:
Artificial Intelligence, a Modern Approach
3rd edition

Today's Menu

- Recap
- Informed Search Methods
 - Heuristics
 - Greedy Search
 - A* Tree Search

Recap: uninformed search

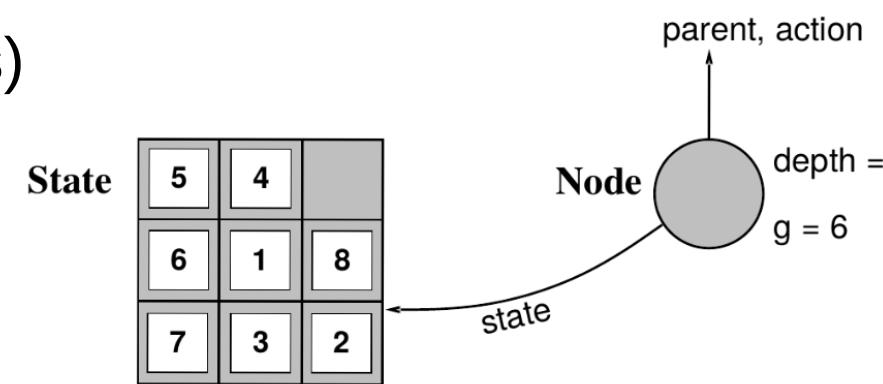
- **Search problem**

- States (configurations of the world)
- Actions and costs
- Successor function (world dynamics)
- Start state and goal test

Recap from last week

- **Search tree**

- Nodes:
represent plans for reaching states
- Plans have costs
(sum of action costs)



- **Search algorithm**

- Systematically builds a search tree
- Chooses an ordering of the fringe (unexplored nodes)
- Optimal: finds cheapest plans (aka sequence of actions)

Fringe and *Frontier*
are synonyms

Tree search algorithms

Recap from last week

Basic idea:

offline, simulated exploration of state space
by generating successors of already-explored states
(a.k.a. **expanding** states)

function TREE-SEARCH(*problem*) **returns** a solution, or failure

initialize the frontier using the initial state of *problem*

loop do

if the frontier is empty **then return** failure

 choose a leaf node and remove it from the frontier

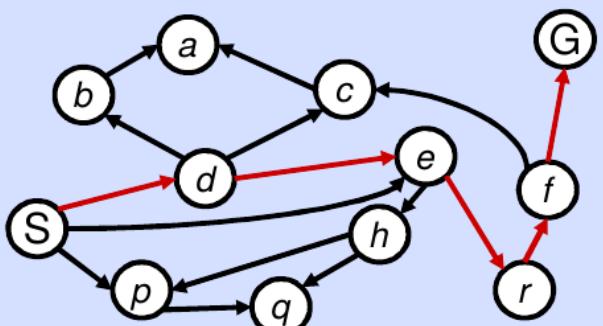
if the node contains a goal state **then return** the corresponding solution

 expand the chosen node and add the resulting nodes to the frontier

end

State Graphs vs. Search Trees

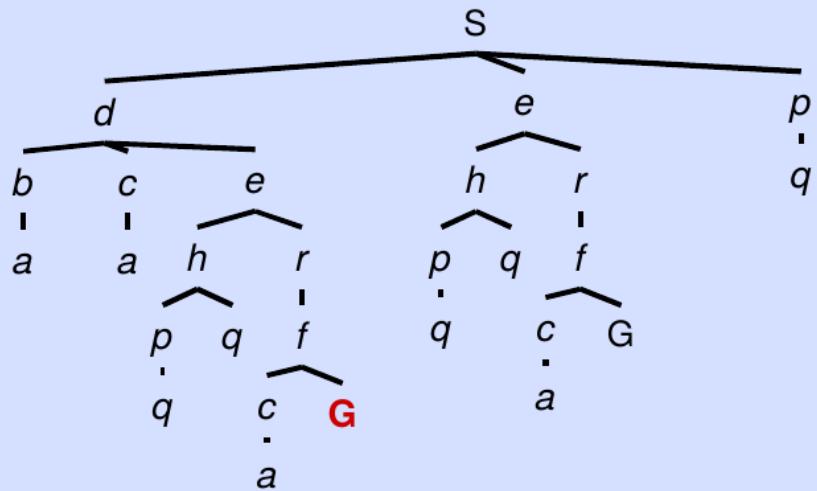
State Graph



*Each NODE in in
the search tree is
an entire PATH in
the problem graph.*

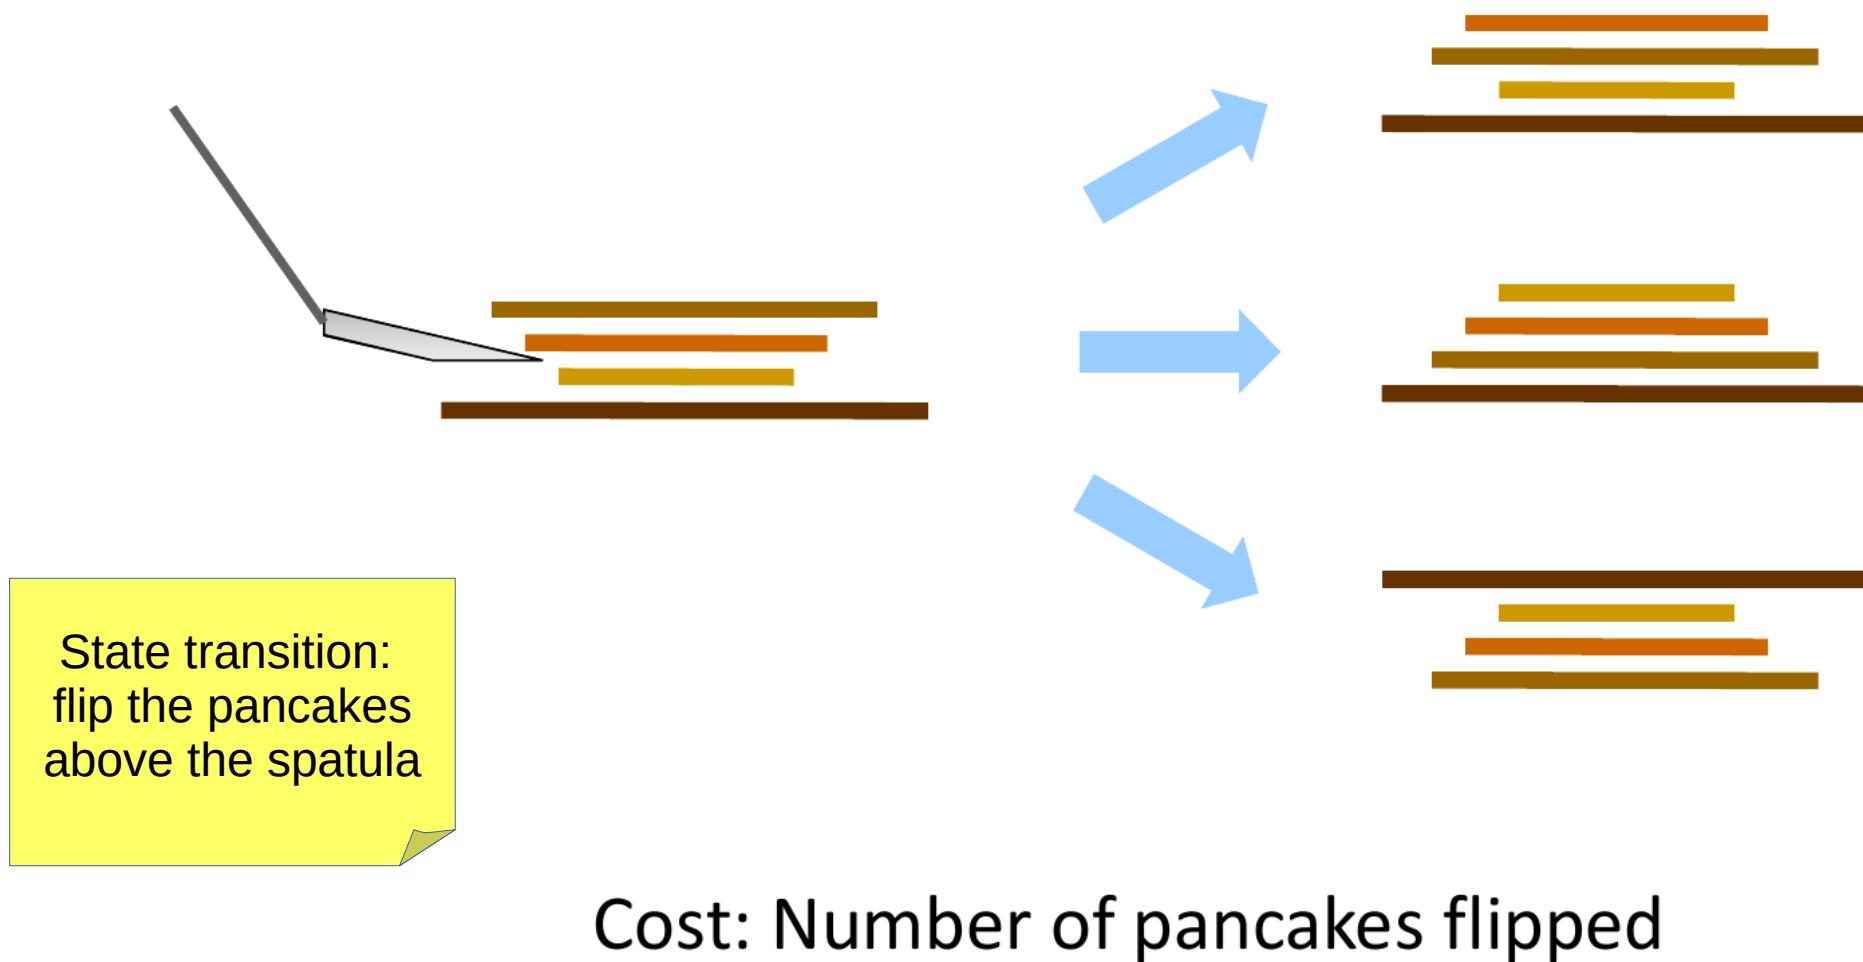
*We construct both
on demand – and
we construct as
little as possible.*

Search Tree



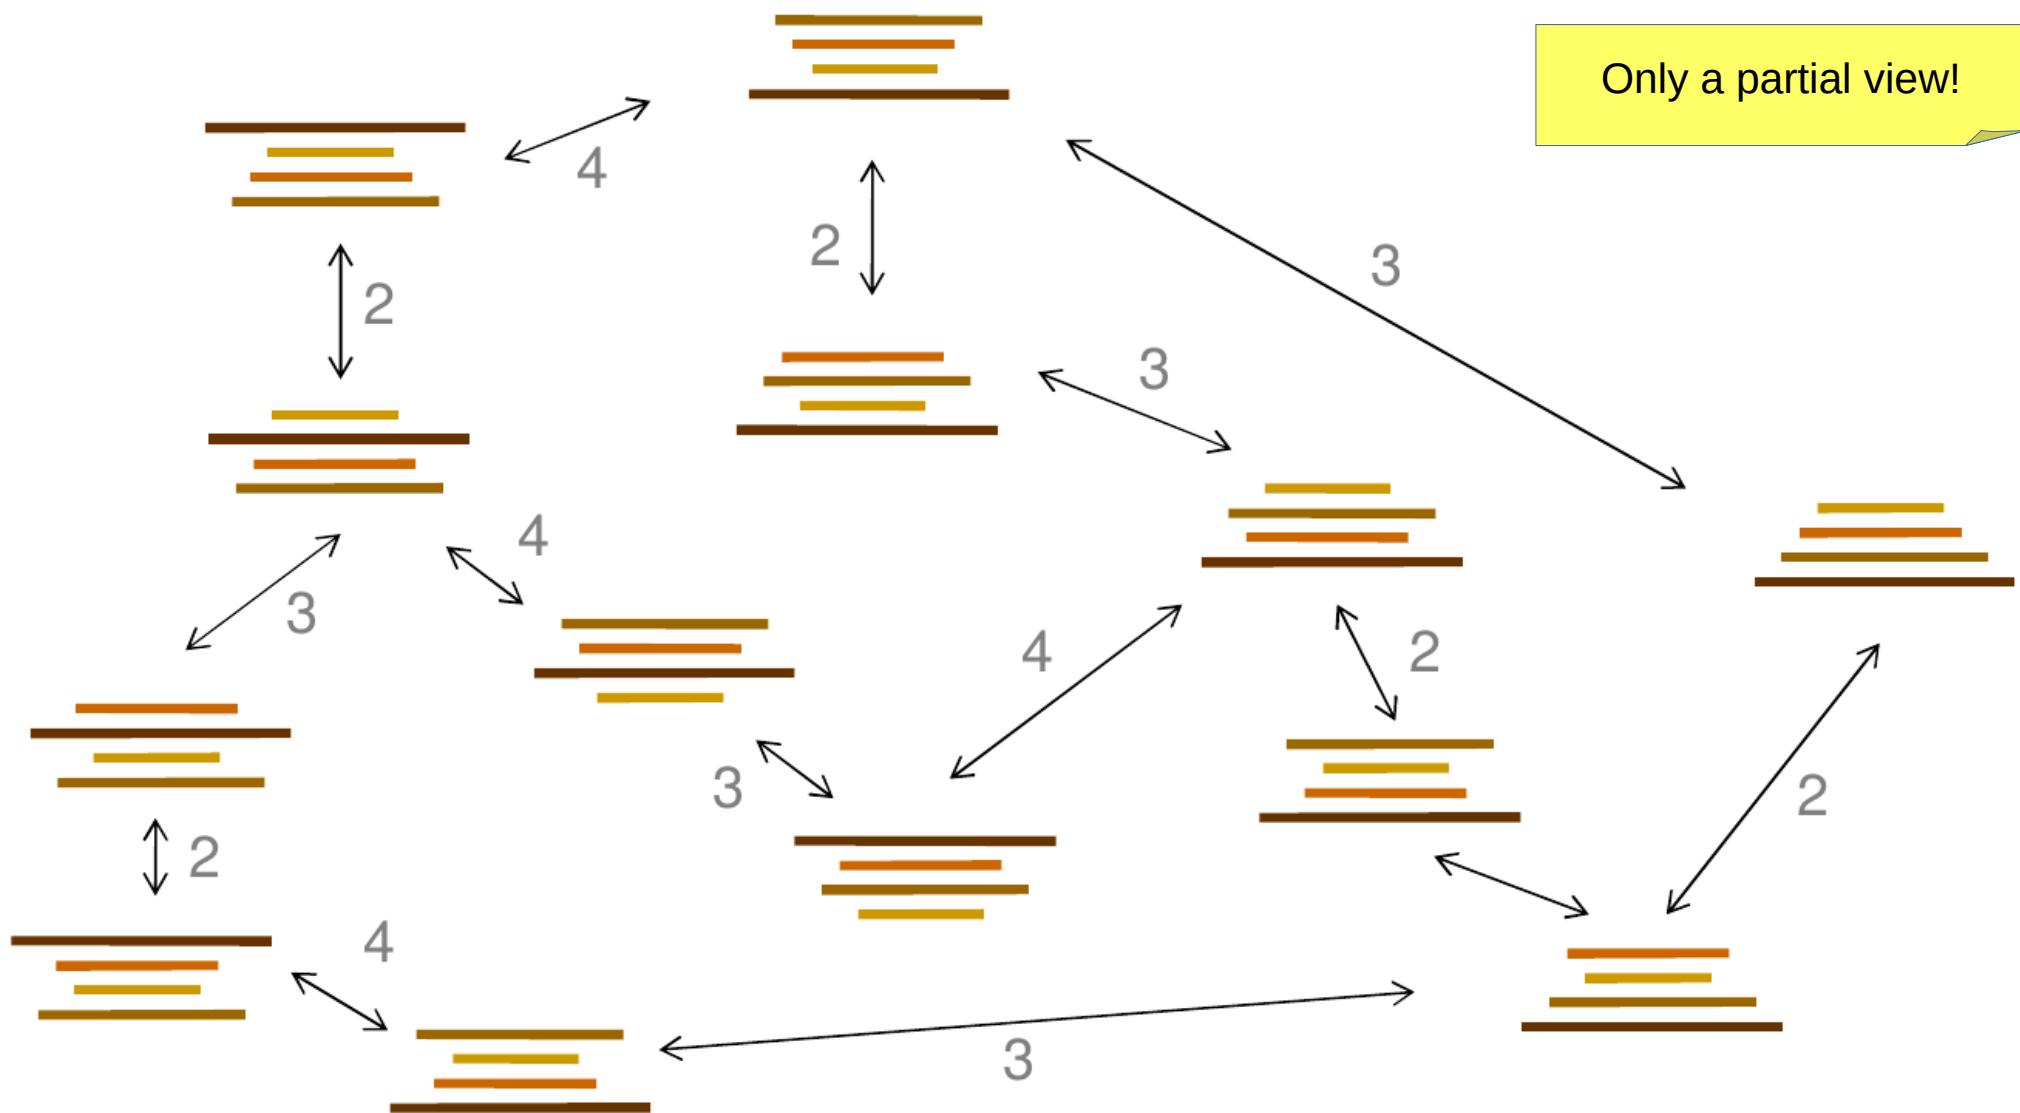
Don't confuse these two!

Example: pancake problem



Example: pancake problem

State space graph with costs as weights



Example: pancake problem

```
function TREE-SEARCH( problem, strategy) returns a solution, or failure
```

```
    initialize the search tree using the initial state of problem
```

```
loop do
```

```
    if there are no candidates for expansion then return failure
```

```
    choose a leaf node for expansion according to strategy
```

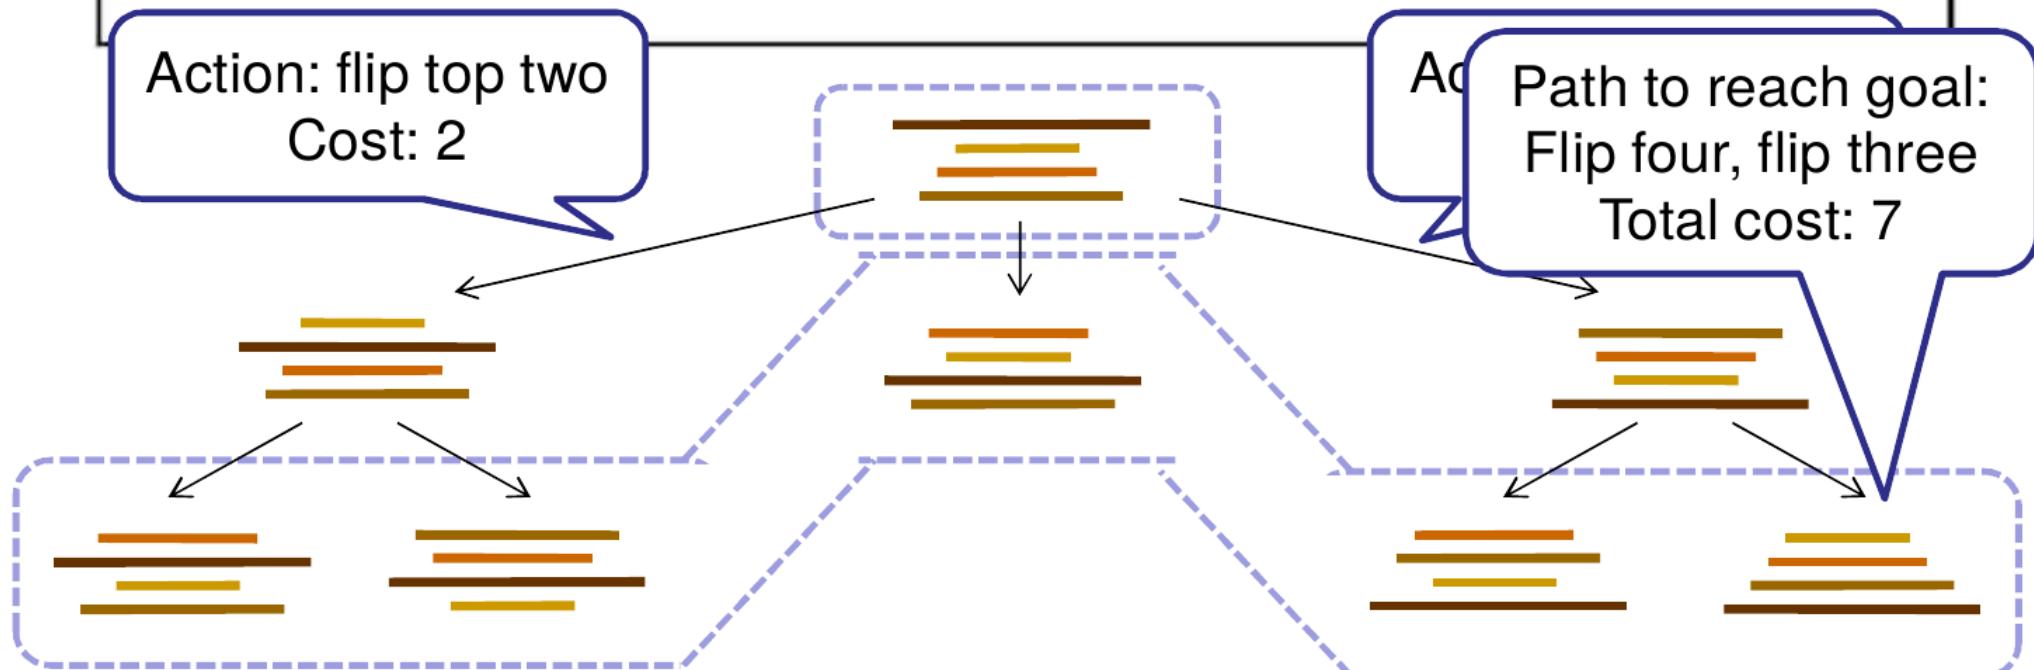
```
    if the node contains a goal state then return the corresponding solution
```

```
    else expand the node and add the resulting nodes to the search tree
```

```
end
```

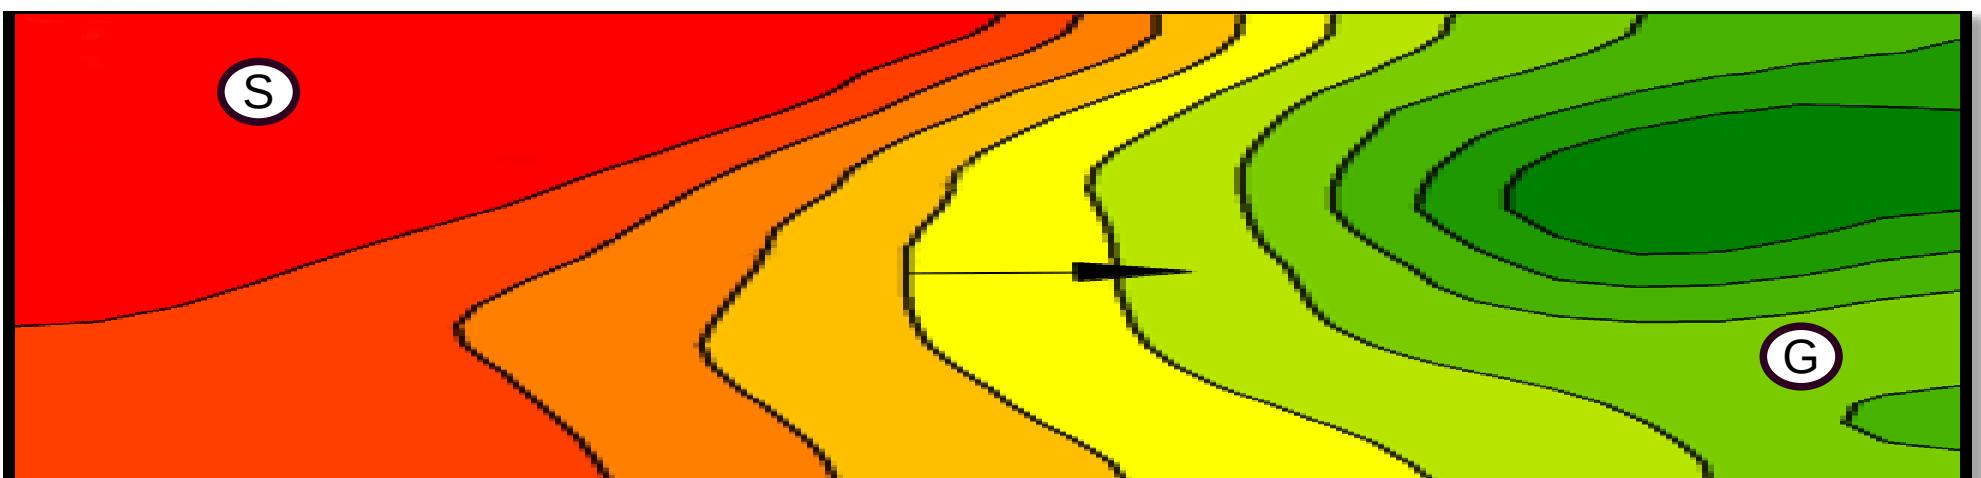
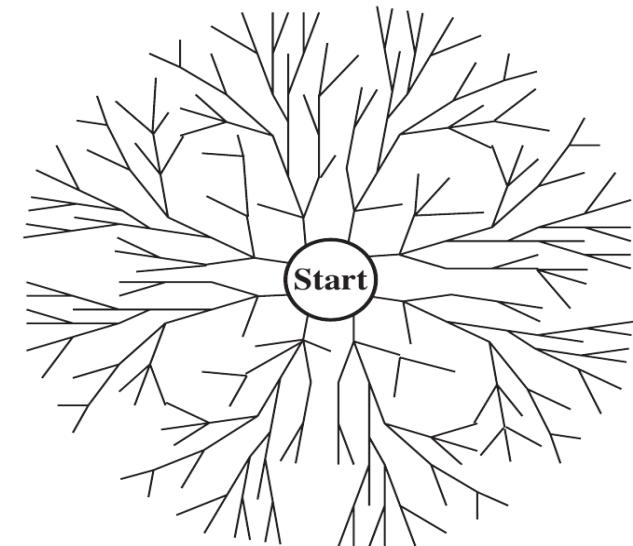
Action: flip top two
Cost: 2

Action: Path to reach goal:
Flip four, flip three
Total cost: 7

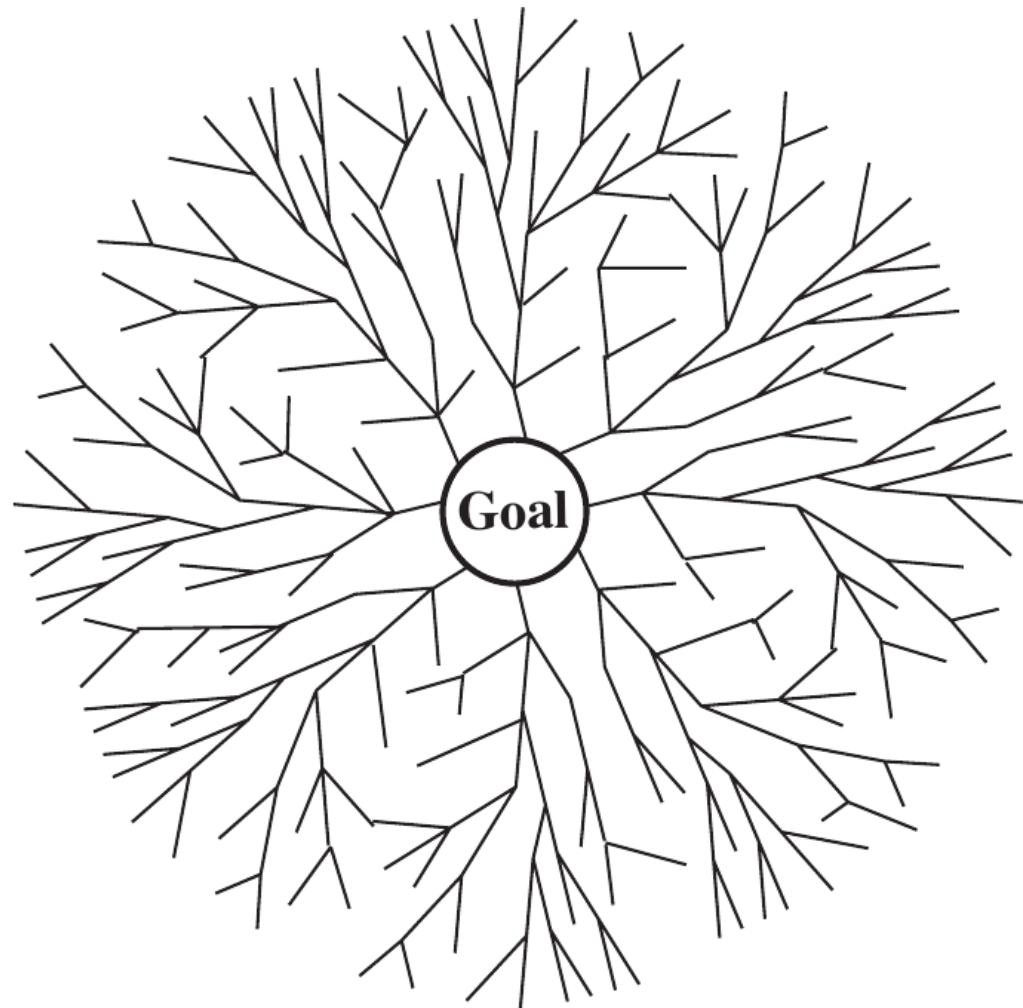
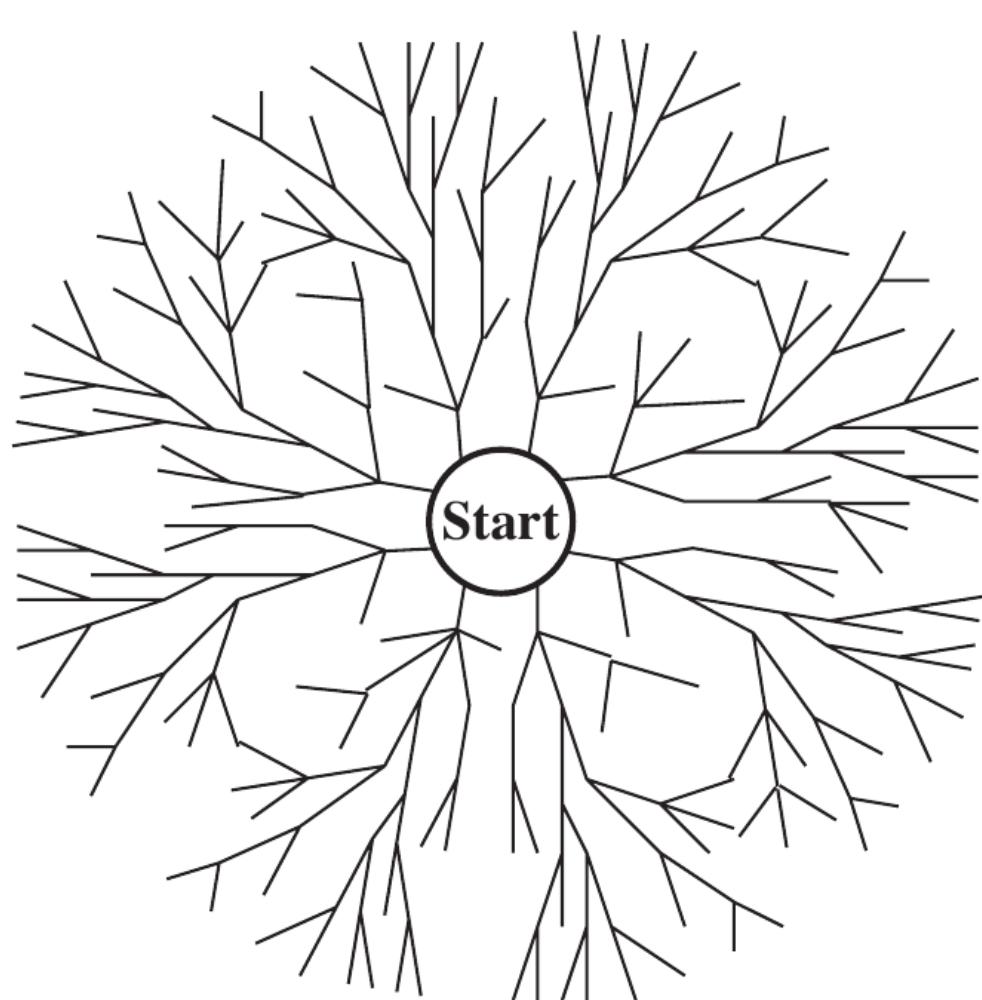


Uniform Cost Search - Issues

- UCS explores increasing cost contours
- The good news:
 - UCS is complete and optimal!
- The bad news:
 - Explores options in every “direction”
 - No information about goal location

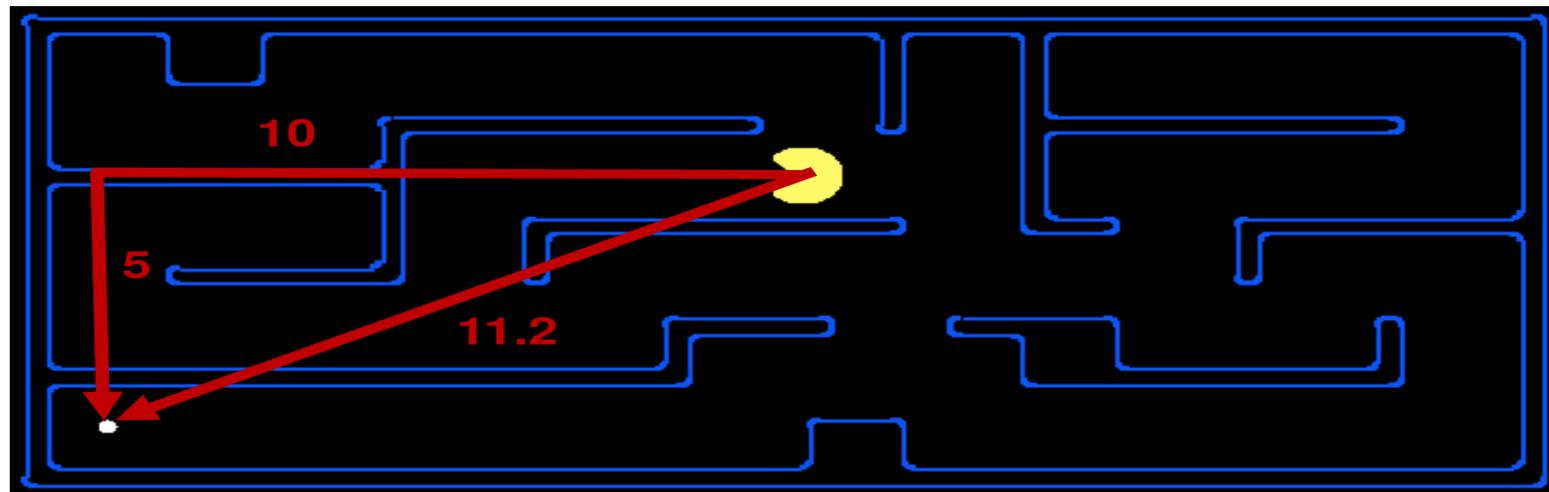


A schematic view of a bidirectional search



Search Heuristics

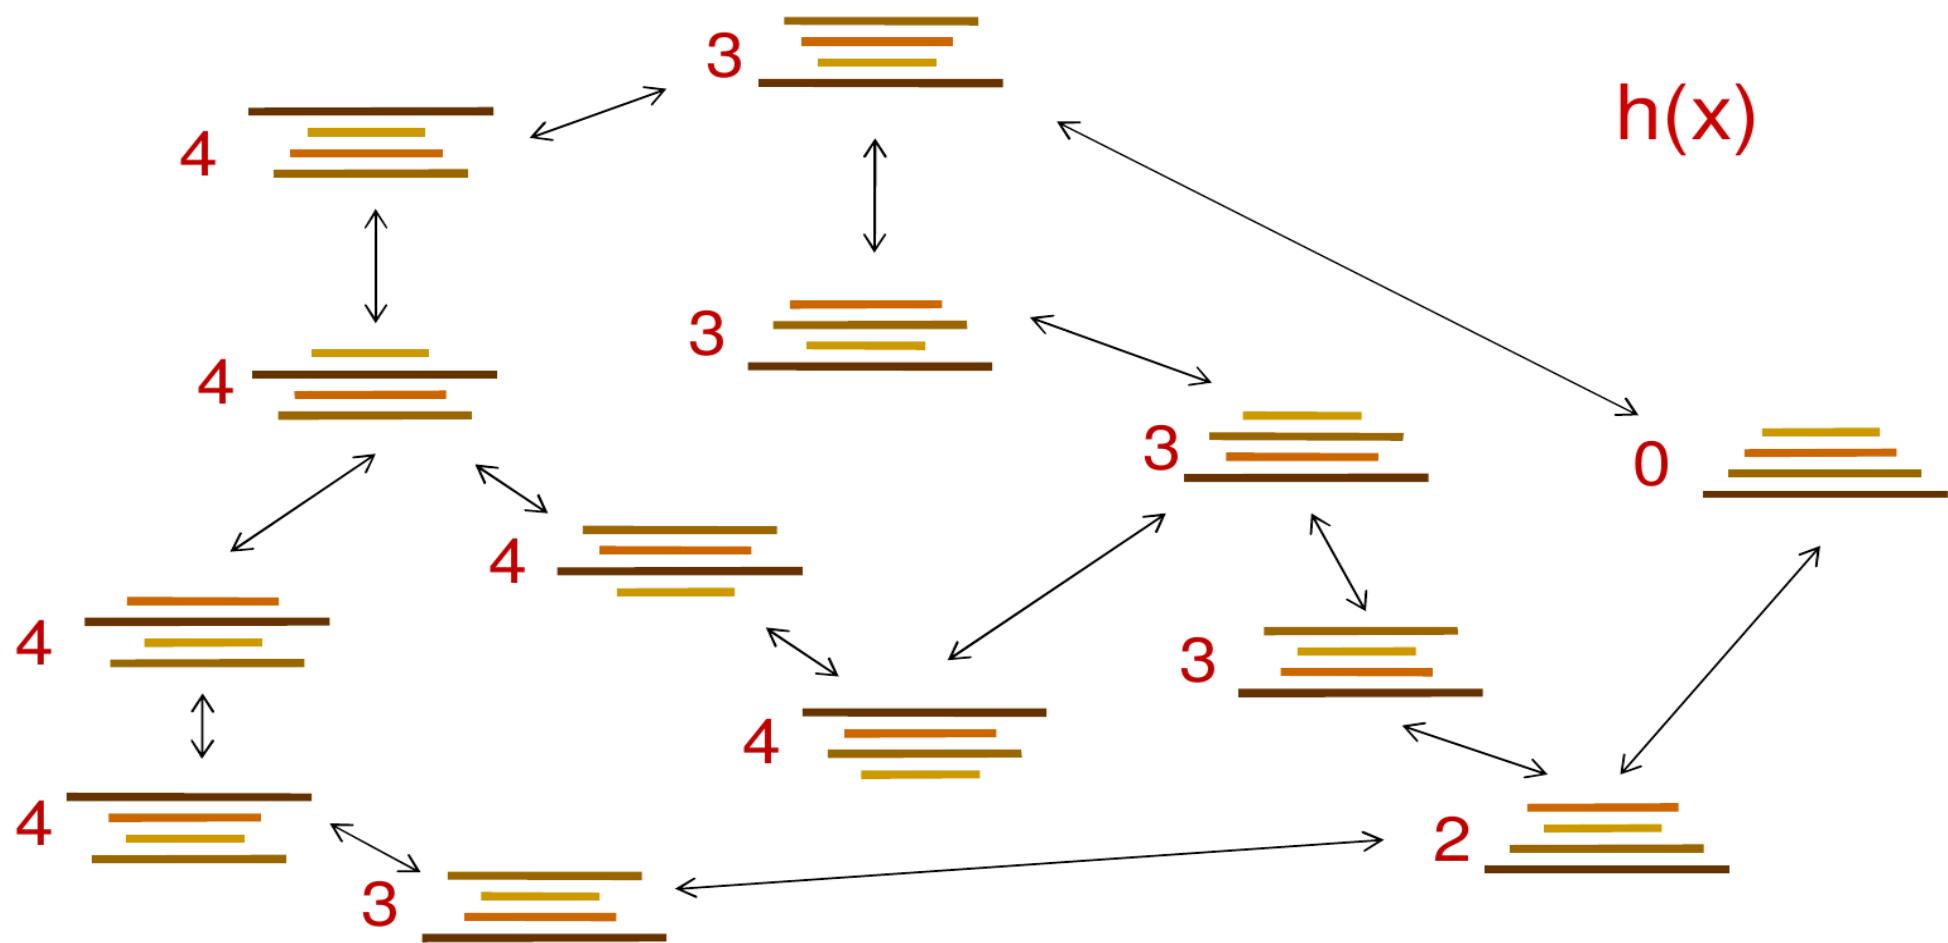
- A **heuristic** is:
 - A function that estimates how close a state is to a goal
 - Designed for a particular search problem
- Examples: **Manhattan distance**, **Euclidean distance** for shortest path finding





Example: Heuristic Function

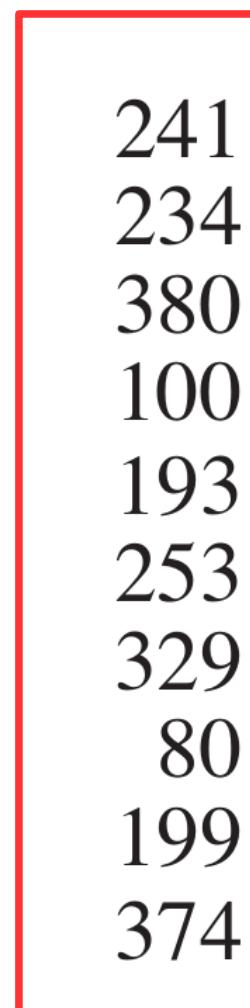
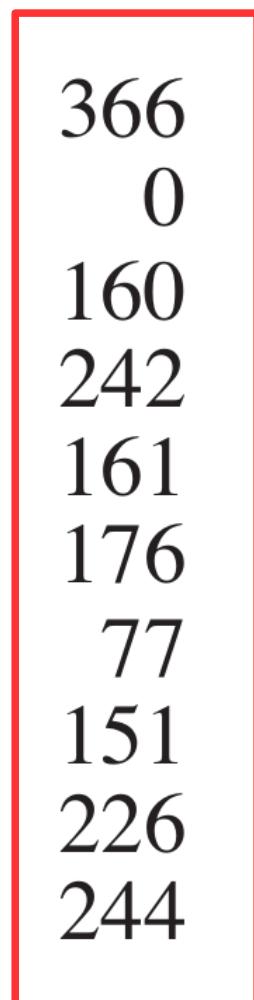
- Heuristic: the index of the largest pancake that is still out of place (1 smallest, 4 largest)



Example Heuristic: Straight-line distances to Bucharest

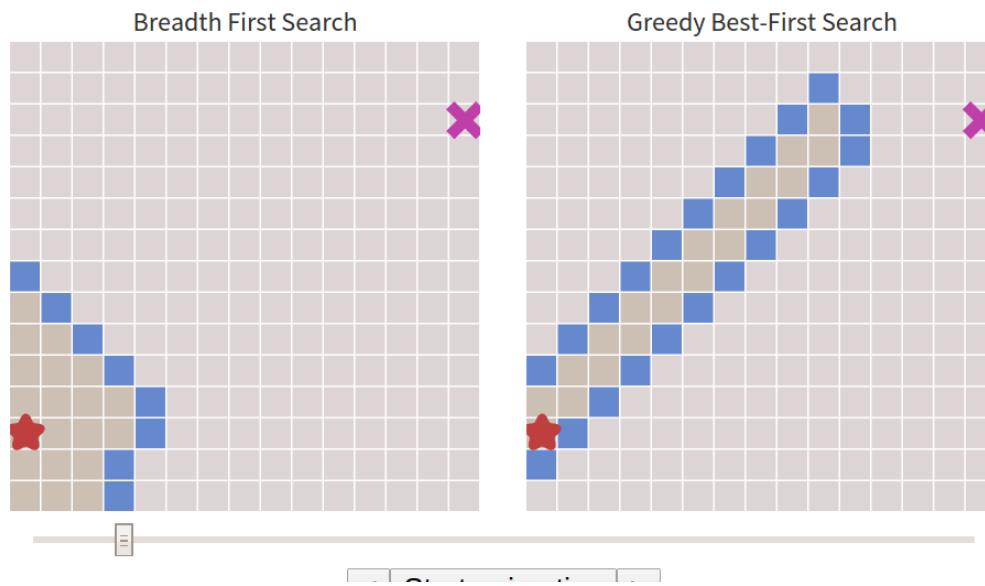
Arad	366
Bucharest	0
Craiova	160
Drobeta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244

Mehadia	241
Neamt	234
Oradea	380
Pitesti	100
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



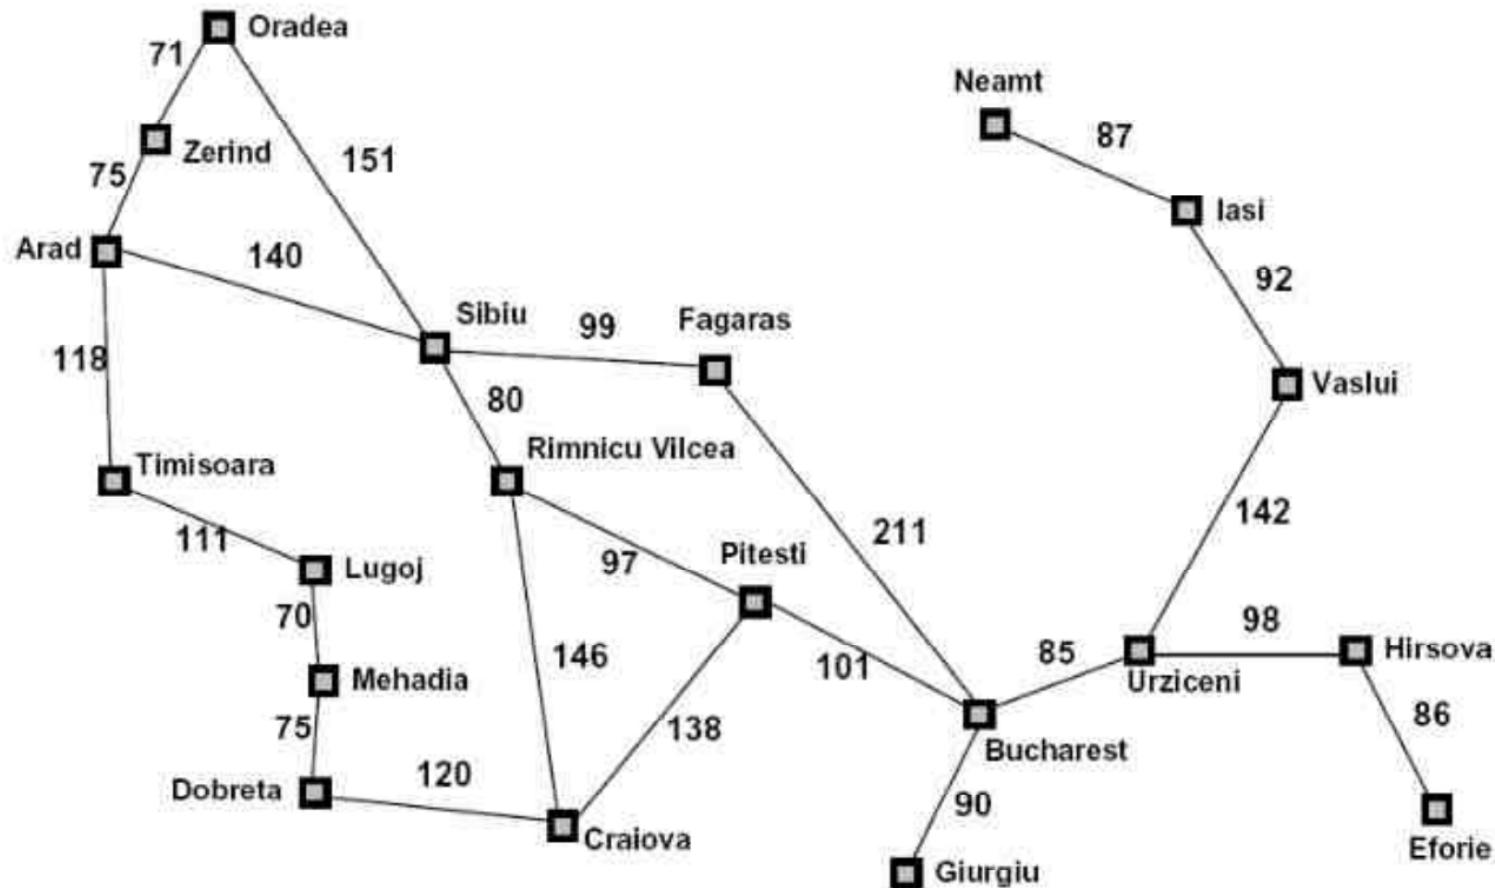
Greedy Search

- Strategy: expand a node that you think is closest to a goal state
 - Heuristic: estimate of distance to nearest goal for each state
- A case where it works well



Greedy search: expand the node that seems the closest to a goal

smallest $h(x)$



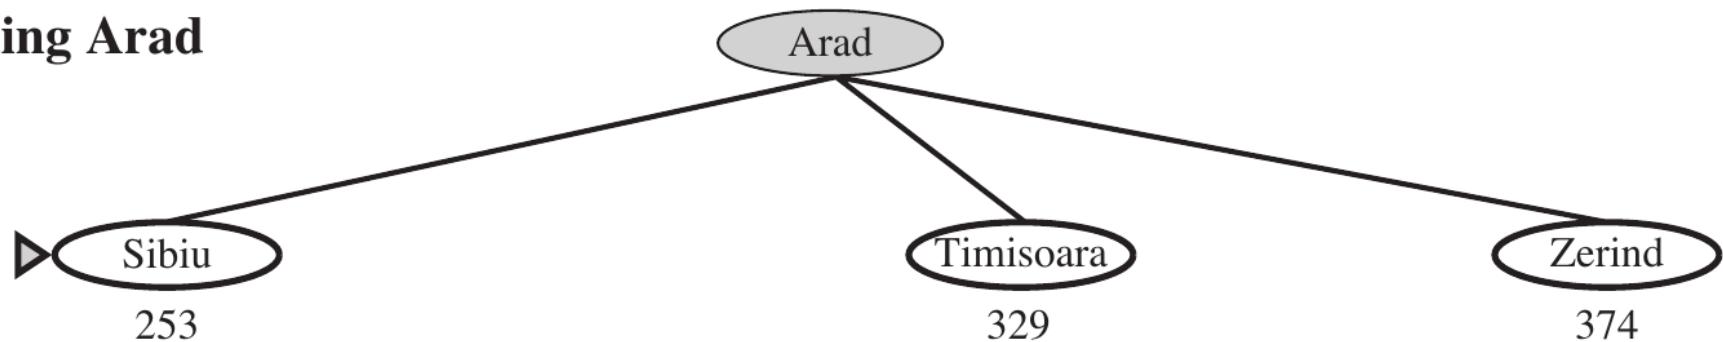
$h(x)$

Greedy best-first tree search for Bucharest with the straight-line distance heuristic

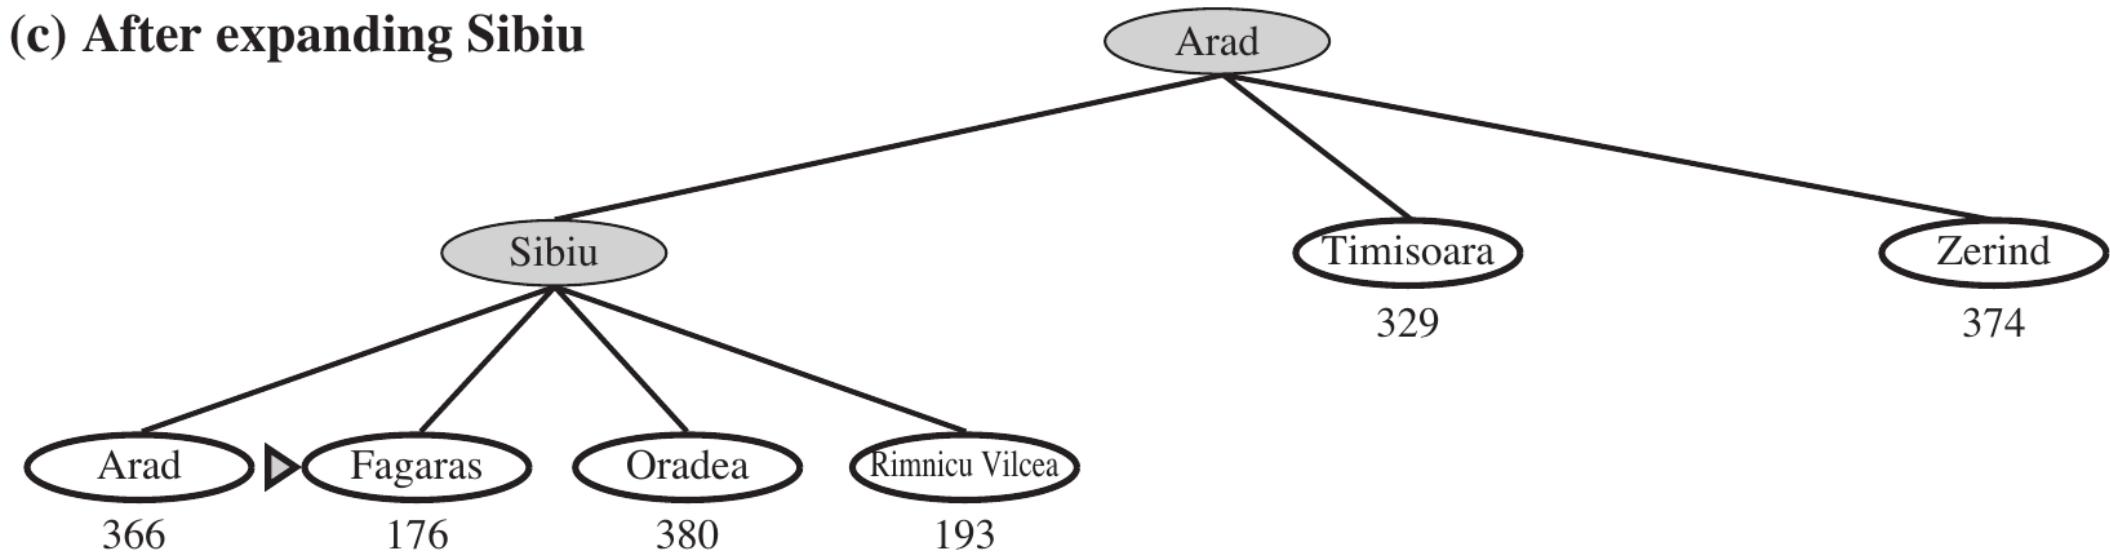
(a) The initial state



(b) After expanding Arad

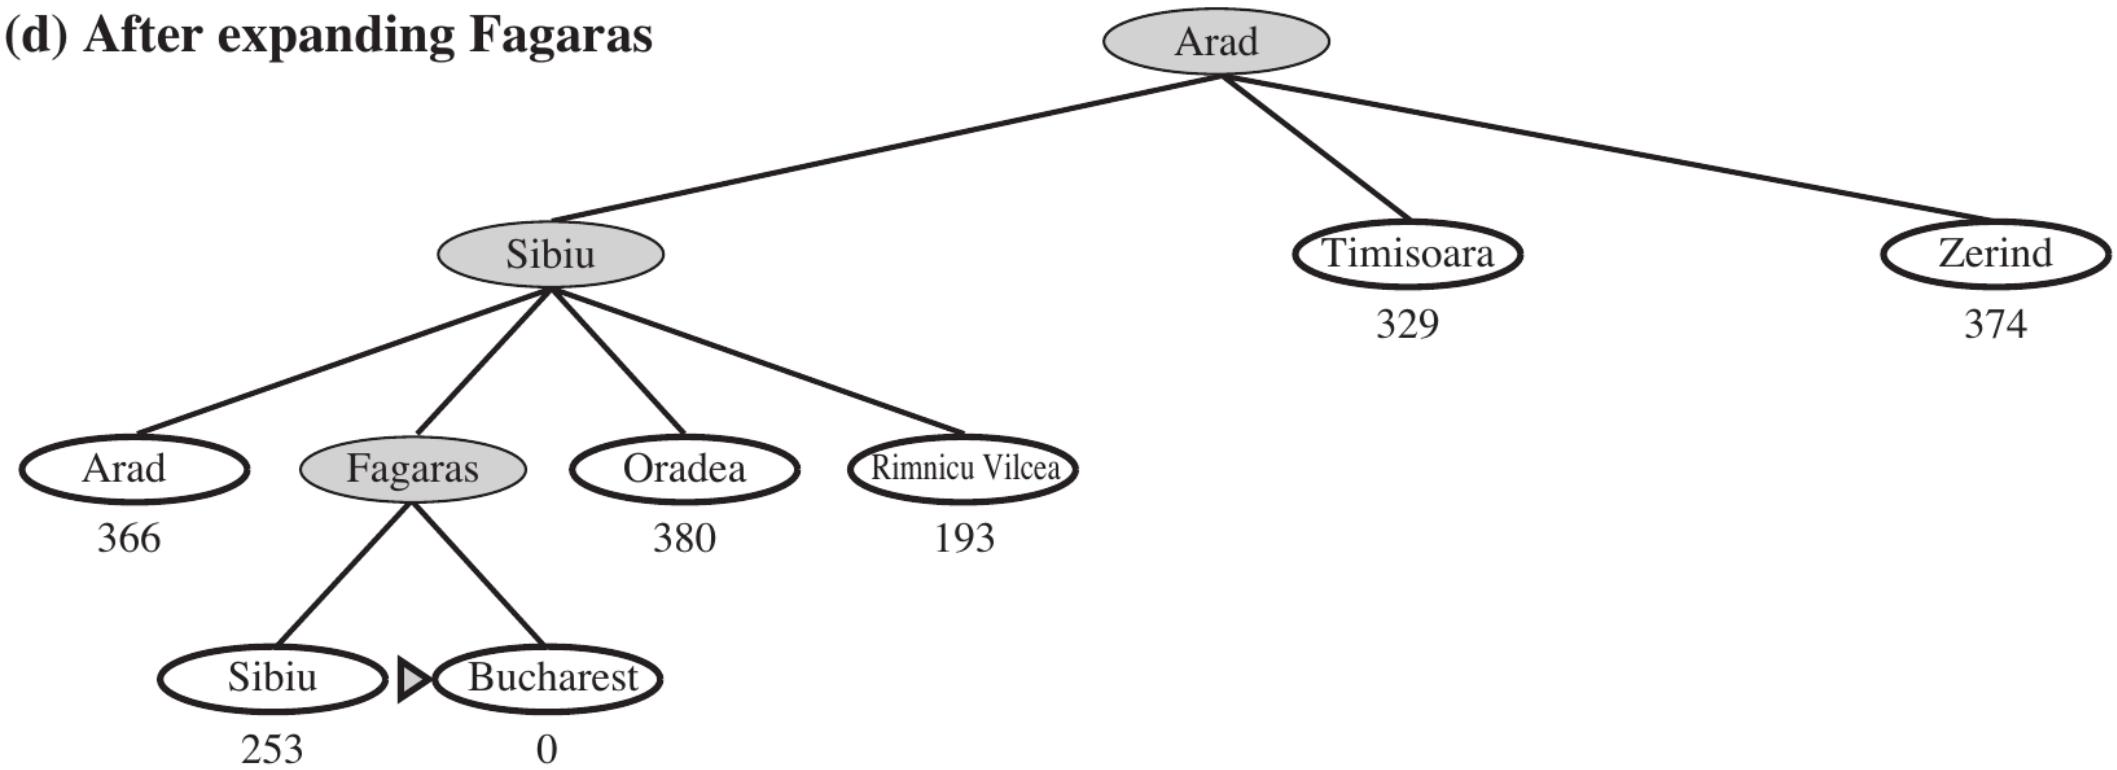


(c) After expanding Sibiu



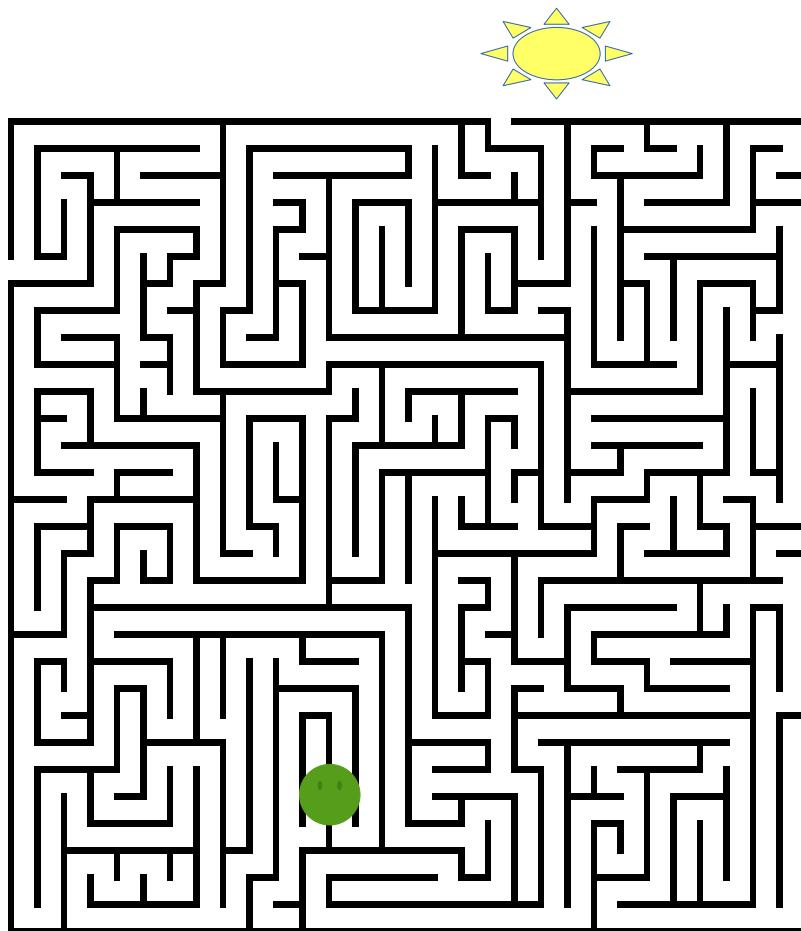
Greedy best-first tree search for Bucharest with the straight-line distance heuristic

(d) After expanding Fagaras

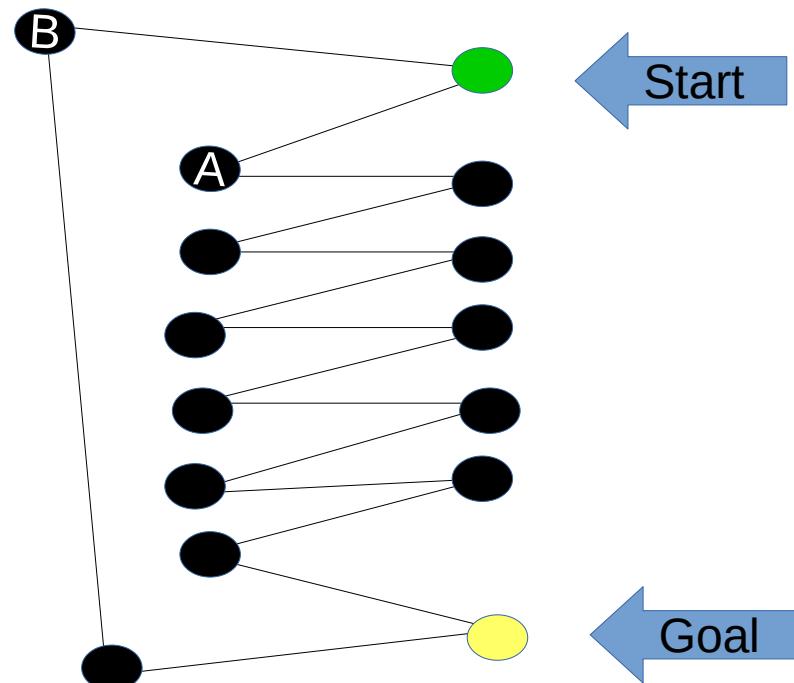


What can go wrong?

Greedy Search



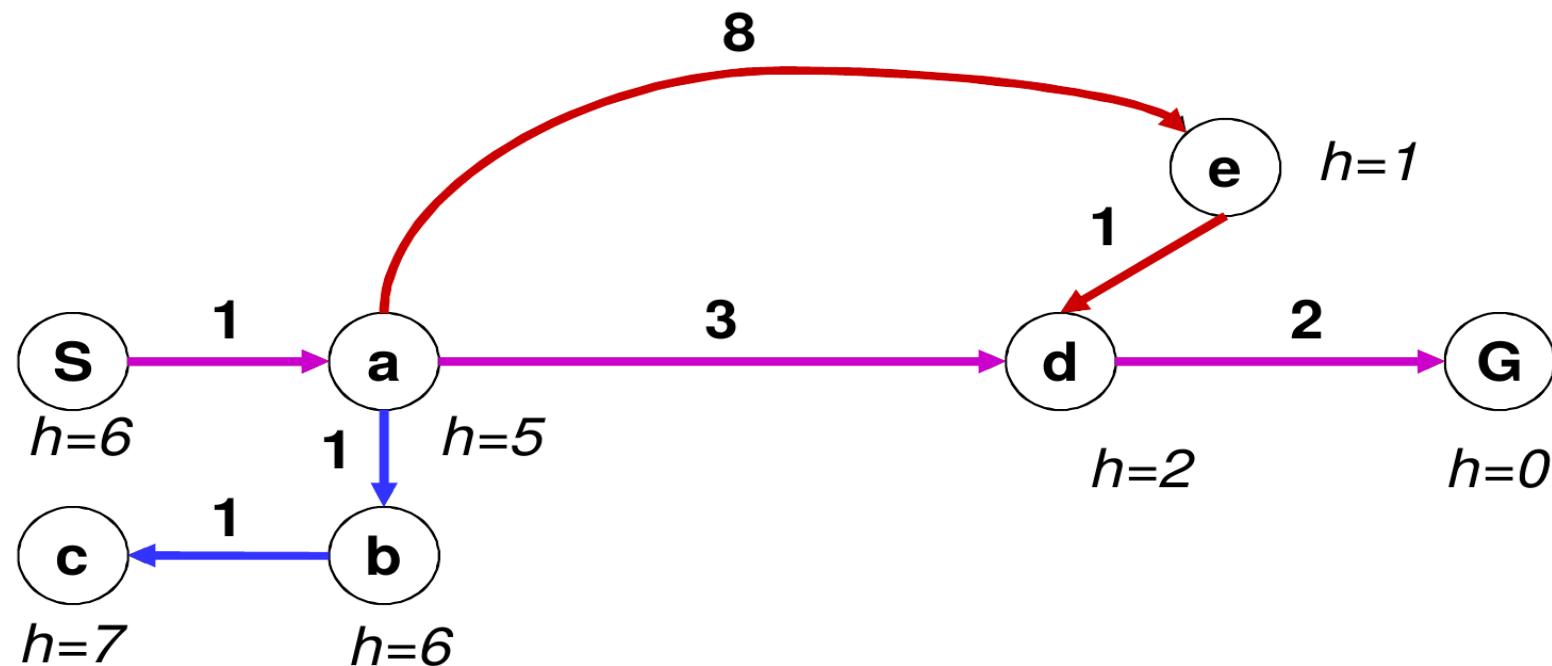
Here $h(x)$ is the straight line distance.



Being greedy (going via A) does not always pay off!
The path via B has a lower cost

Combining UCS and Greedy

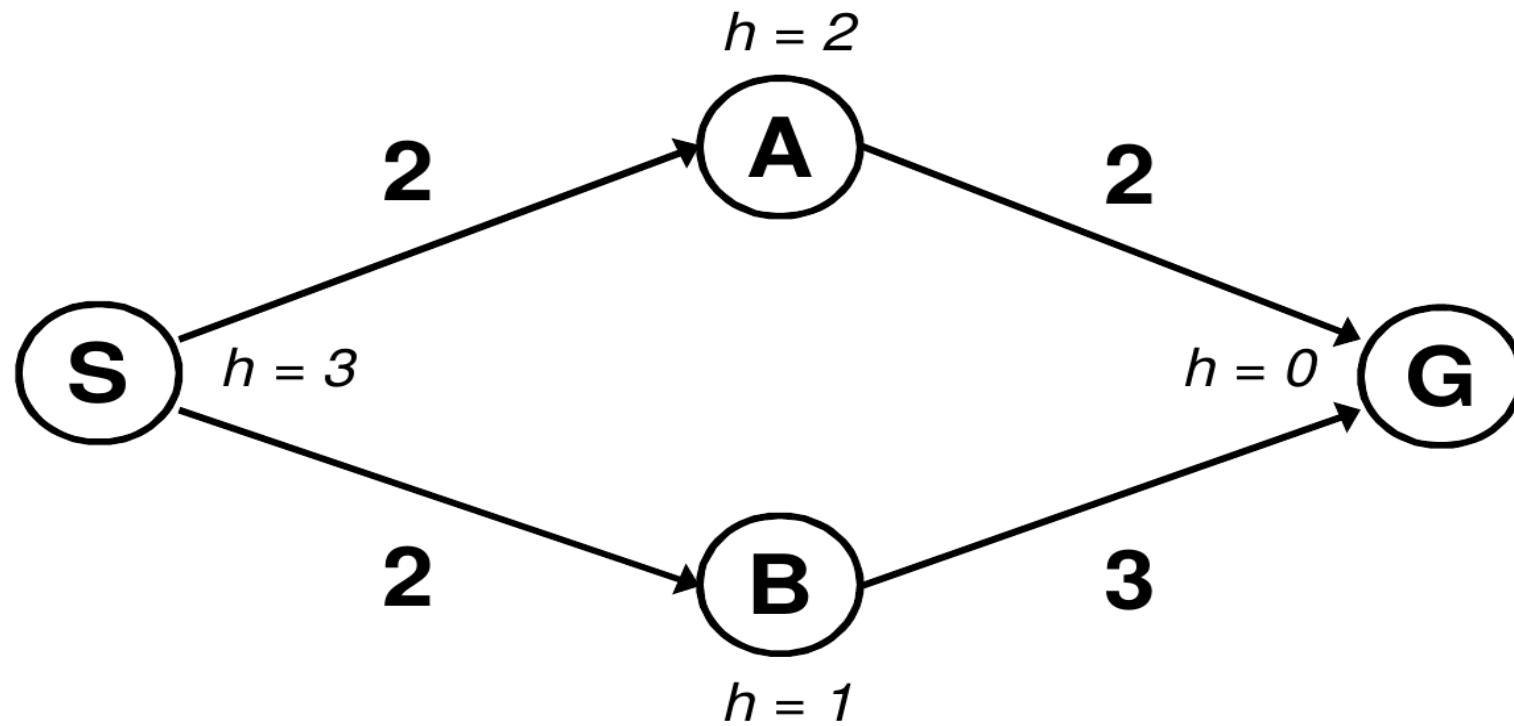
- Uniform-cost orders by path cost, or backward cost $g(n)$
- Greedy orders by goal proximity, or forward cost $h(n)$



A* Search orders by the sum: $f(n) = g(n) + h(n)$

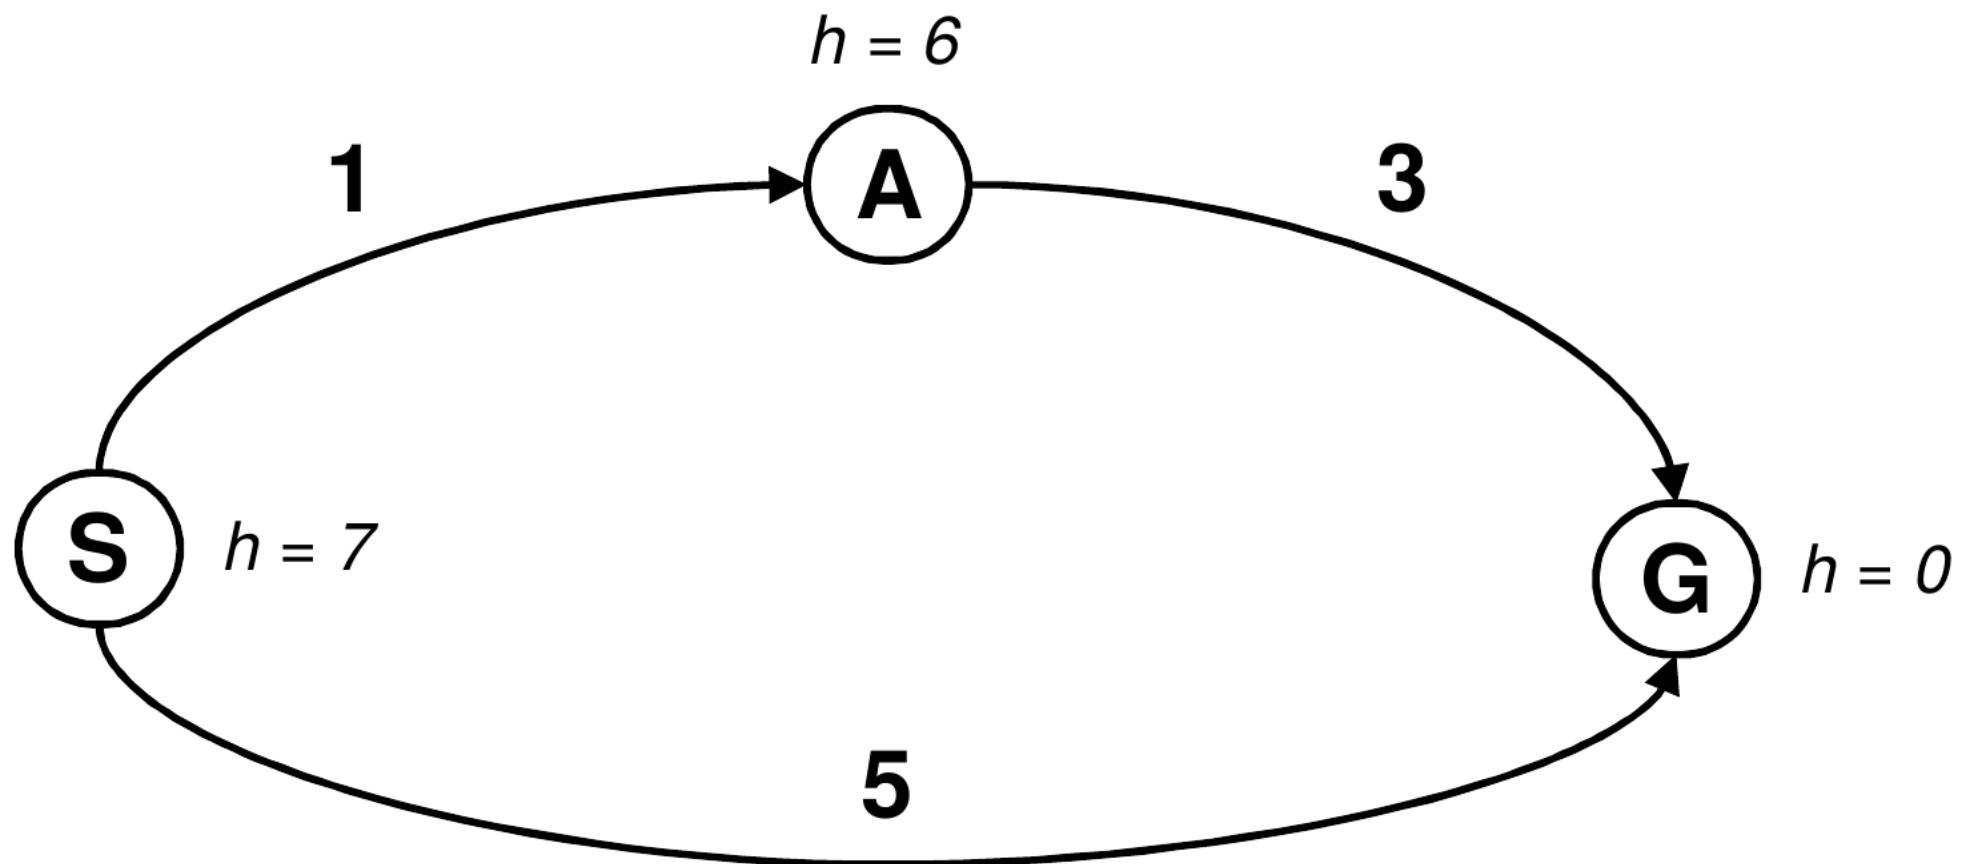
When should A* terminate?

- Should we stop when we enqueue a goal?



- No: only stop when we dequeue a goal

Is A* Optimal?



What went wrong?

Actual bad goal cost < estimated good goal cost

We need estimates to be less than actual costs!

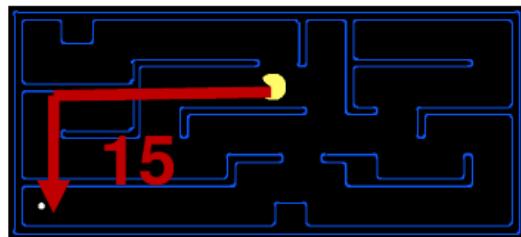
Admissible Heuristics

A heuristic h is *admissible* (optimistic) if:

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal

Examples:



Coming up with admissible heuristics is most of what's involved in using A* in practice.

Optimality of A* Tree Search

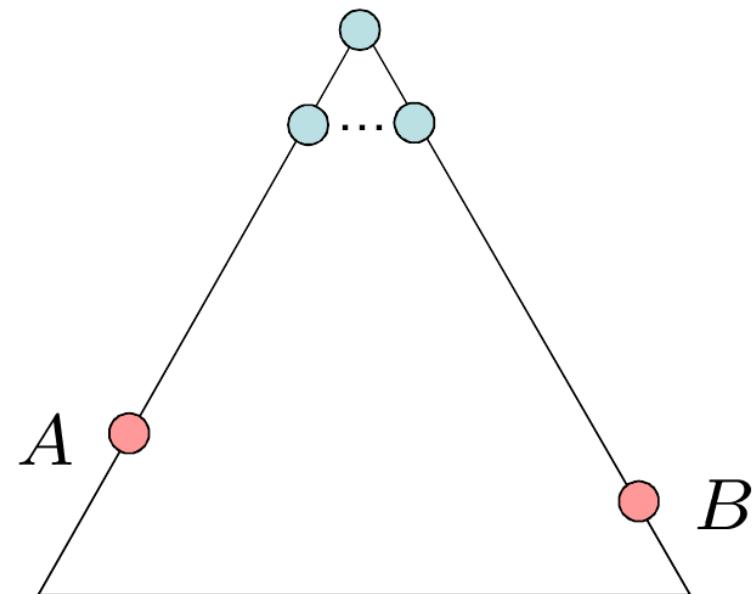
Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim:

- A will exit the fringe before B

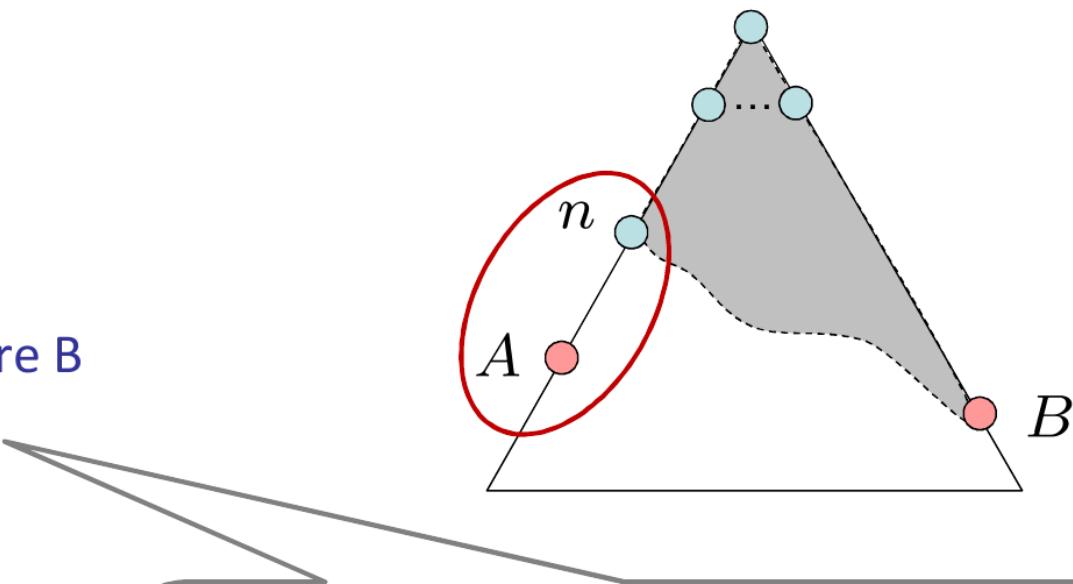
'fringe' and 'frontier' are synonyms



Optimality of A* Tree Search

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$



$$f(n) = g(n) + h(n)$$

$$f(n) \leq g(A)$$

$$g(A) = f(A)$$

Definition of f-cost

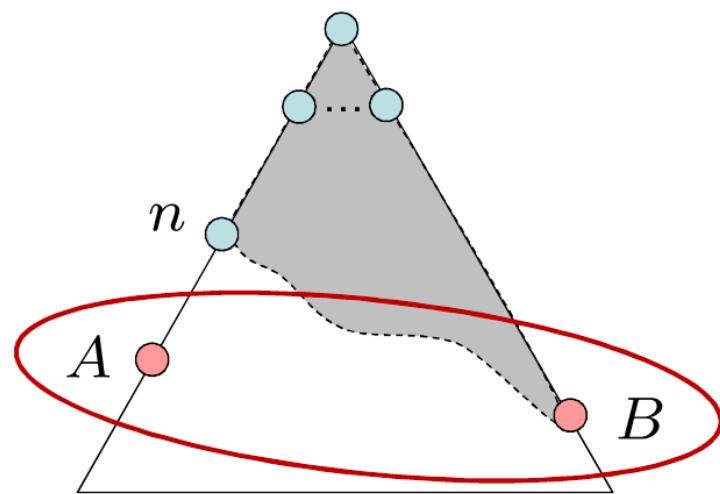
Admissibility of h

$h = 0$ at a goal

Optimality of A* Tree Search

Proof:

- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$



$$g(A) < g(B)$$

$$f(A) < f(B)$$

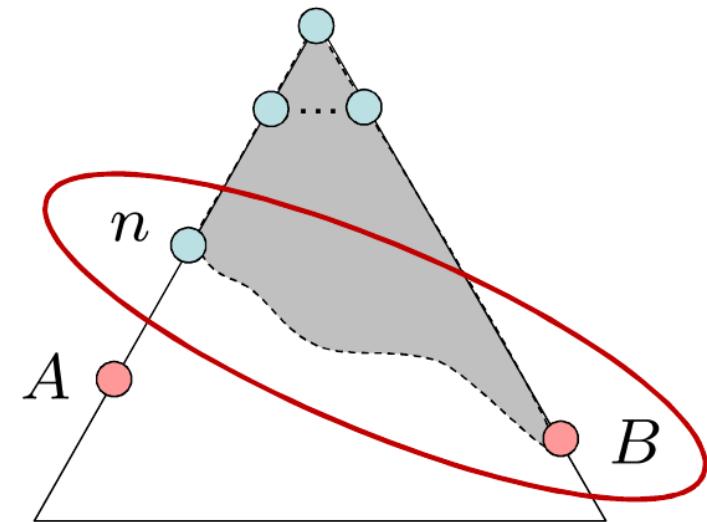
B is suboptimal

$h = 0$ at a goal

Optimality of A* Tree Search

Proof:

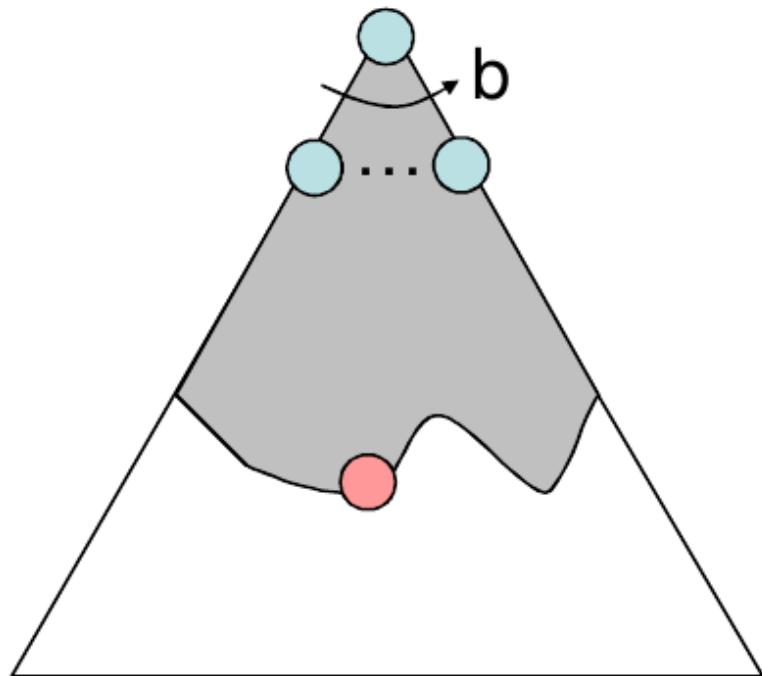
- Imagine B is on the fringe
- Some ancestor n of A is on the fringe, too (maybe A!)
- Claim: n will be expanded before B
 1. $f(n)$ is less or equal to $f(A)$
 2. $f(A)$ is less than $f(B)$
 3. n expands before B
- All ancestors of A expand before B
- A expands before B
- A* search is optimal



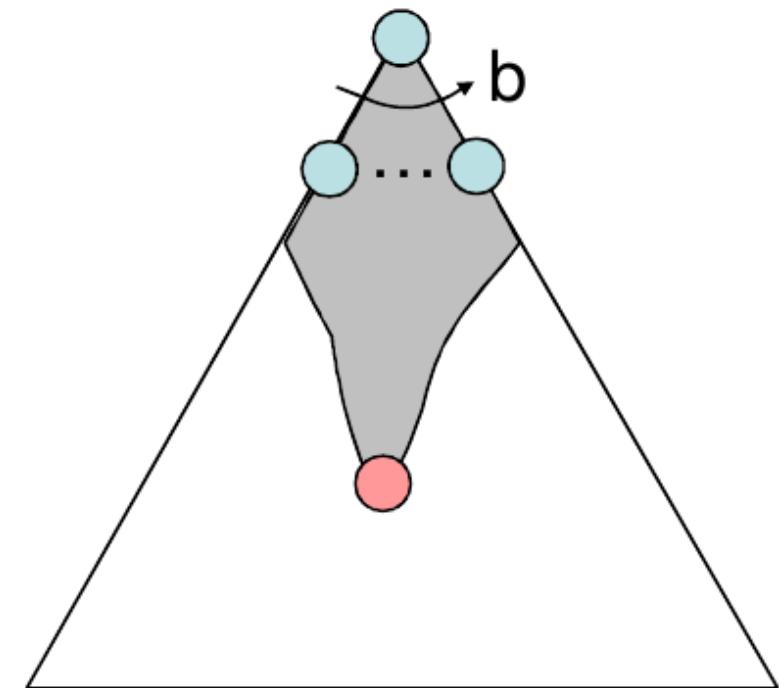
$$f(n) \leq f(A) < f(B)$$

Properties of A*

Uniform-Cost

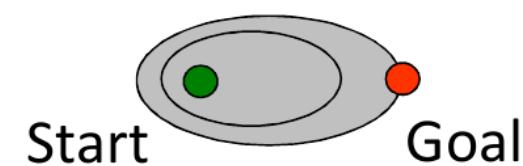
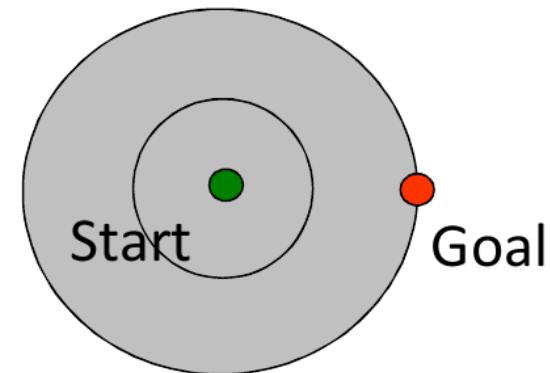


A*



Properties of A*

- Uniform-cost expands equally in all “directions”
- A* expands mainly toward the goal, but does hedge its bets to ensure optimality



Creating Admissible Heuristics

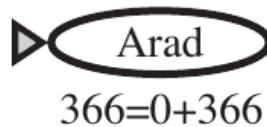
- Most of the work in solving hard search problems optimally is in coming up with **admissible heuristics**
- Often, admissible heuristics are solutions to **relaxed problems**, where new actions are available

Relaxing a problem ~ removing constraints

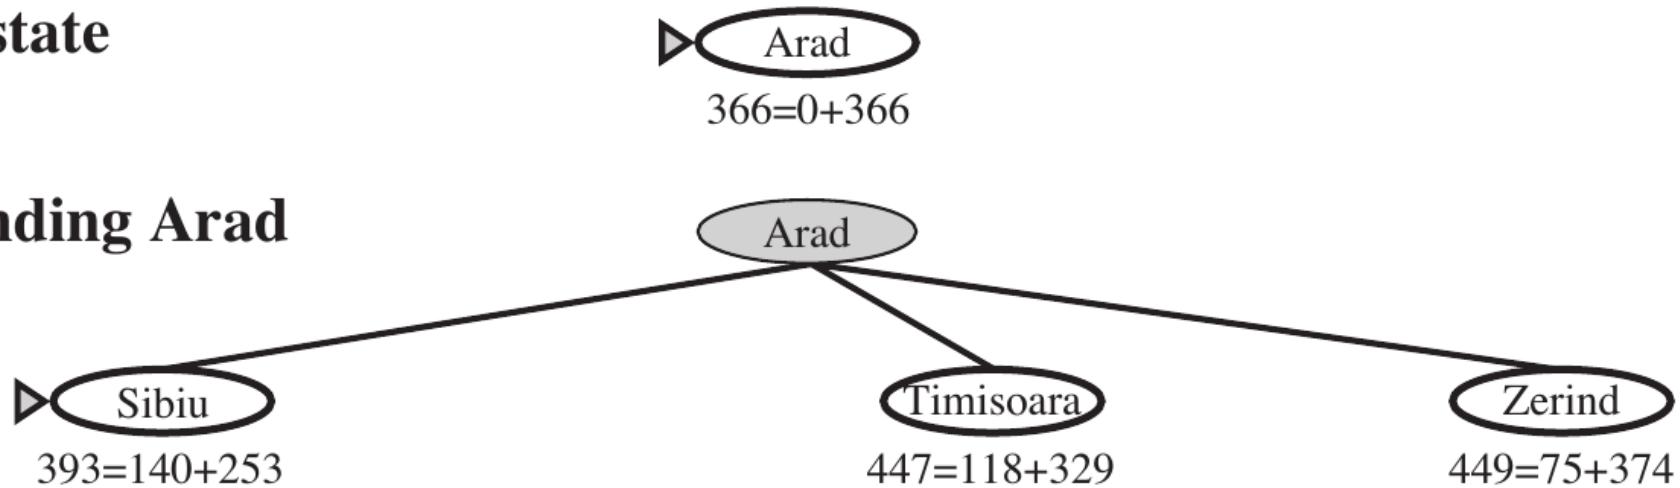
Stages in A* search

Nodes labelled with $f = g+h$

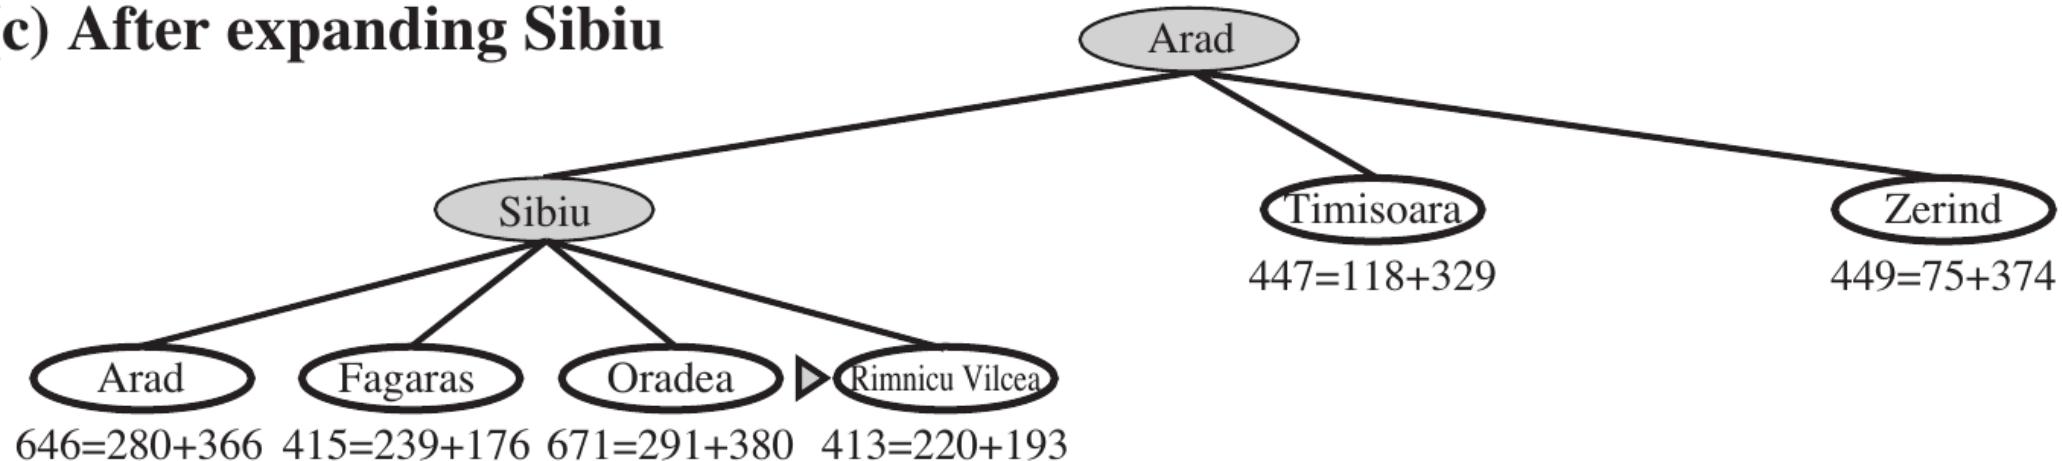
(a) The initial state



(b) After expanding Arad



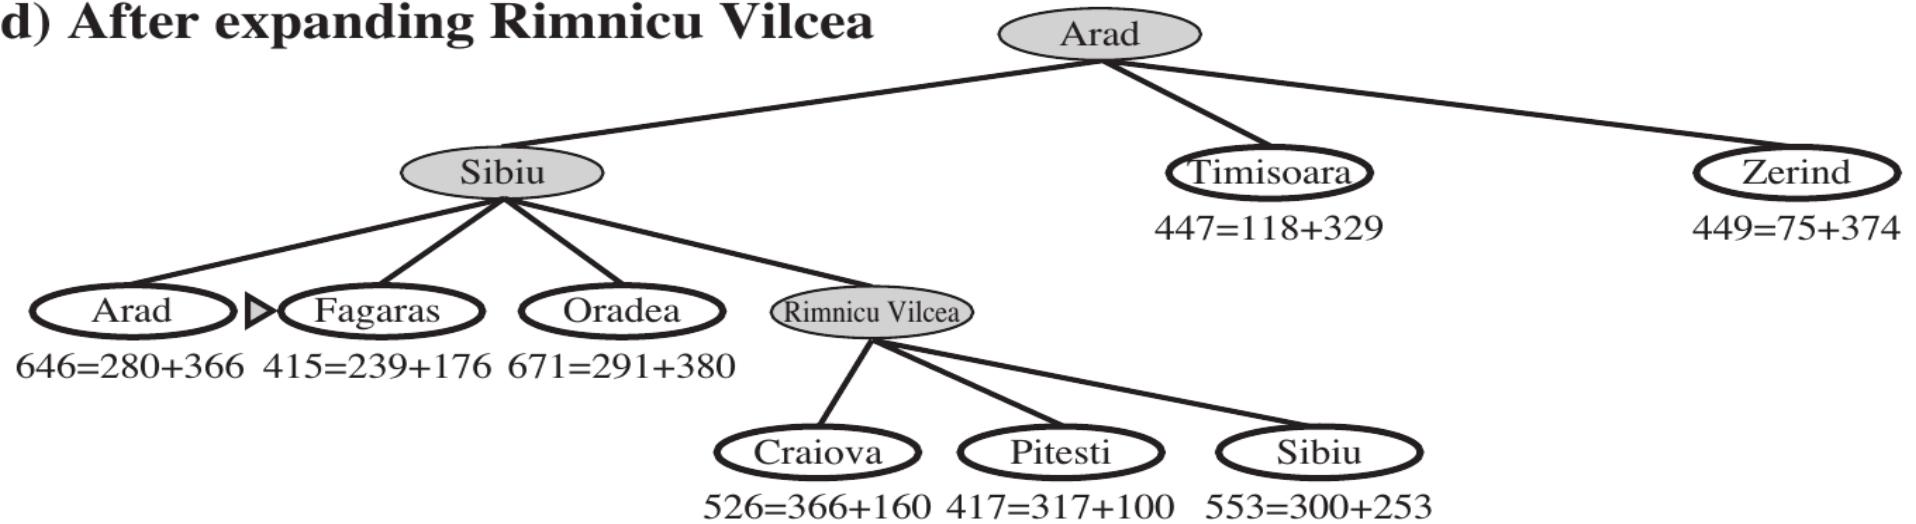
(c) After expanding Sibiu



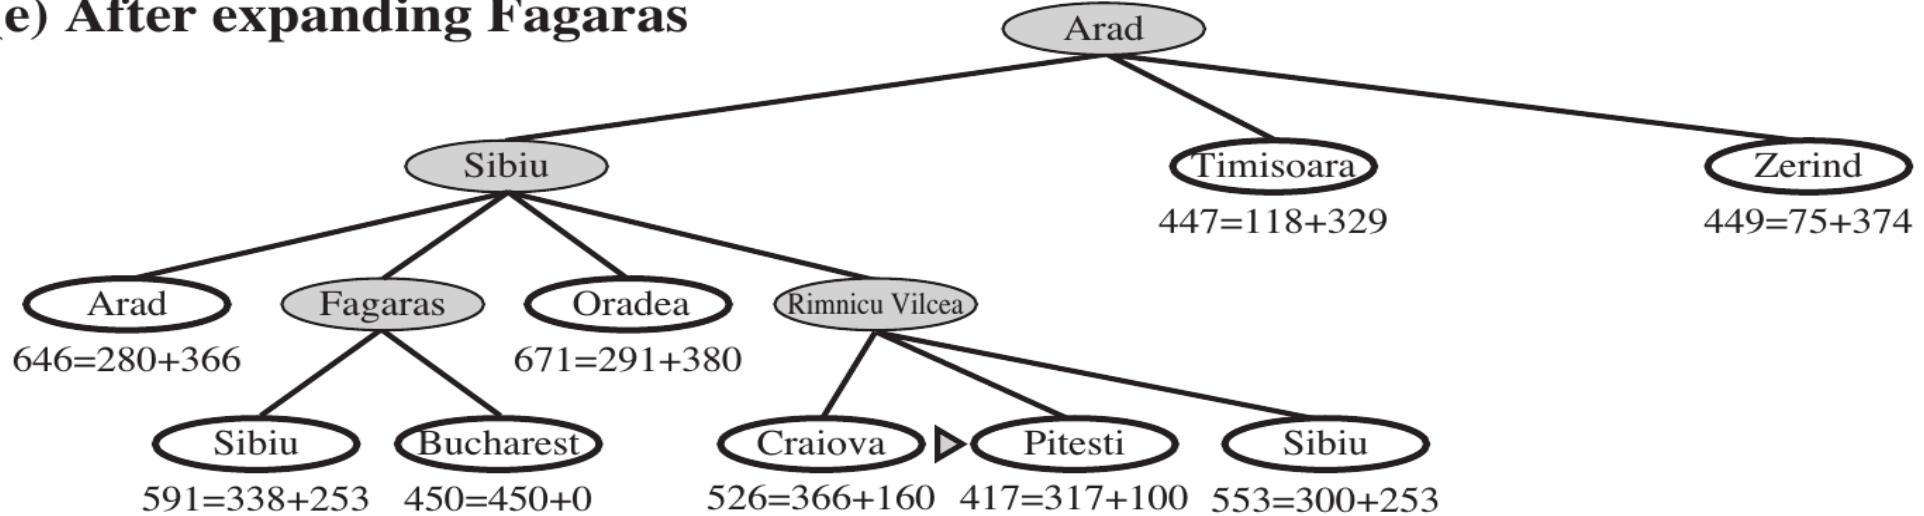
Stages in A* search

Nodes labelled with $f = g+h$

(d) After expanding Rimnicu Vilcea



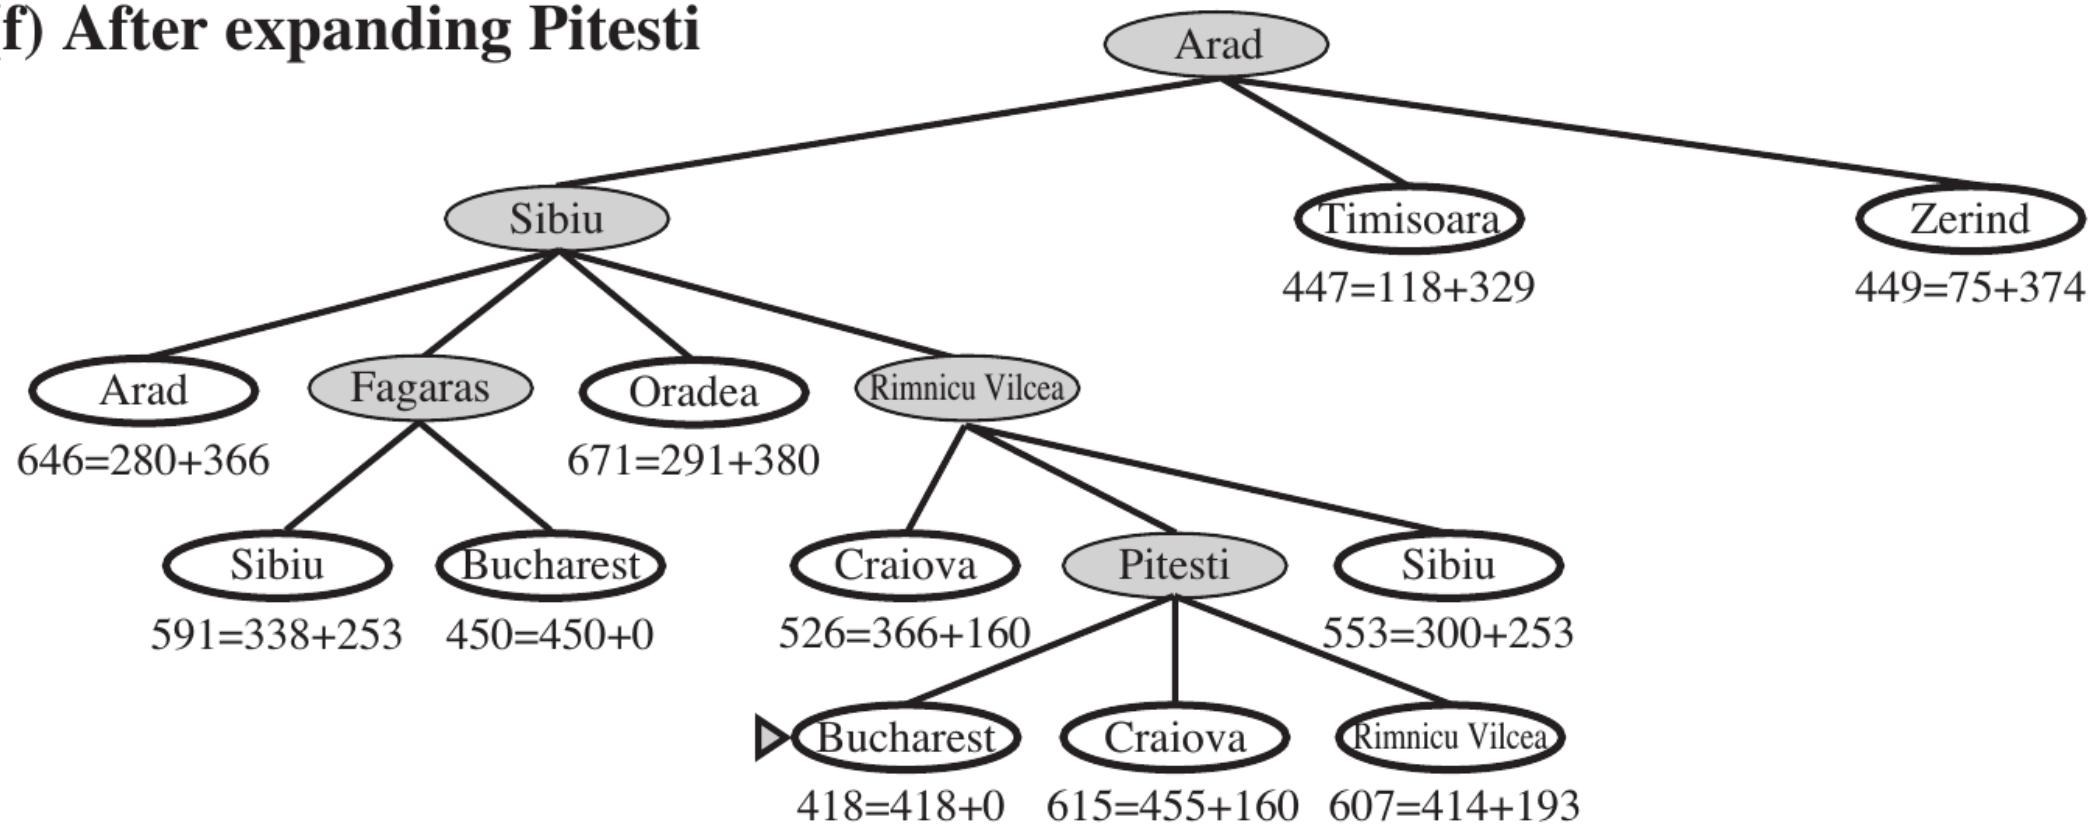
(e) After expanding Fagaras



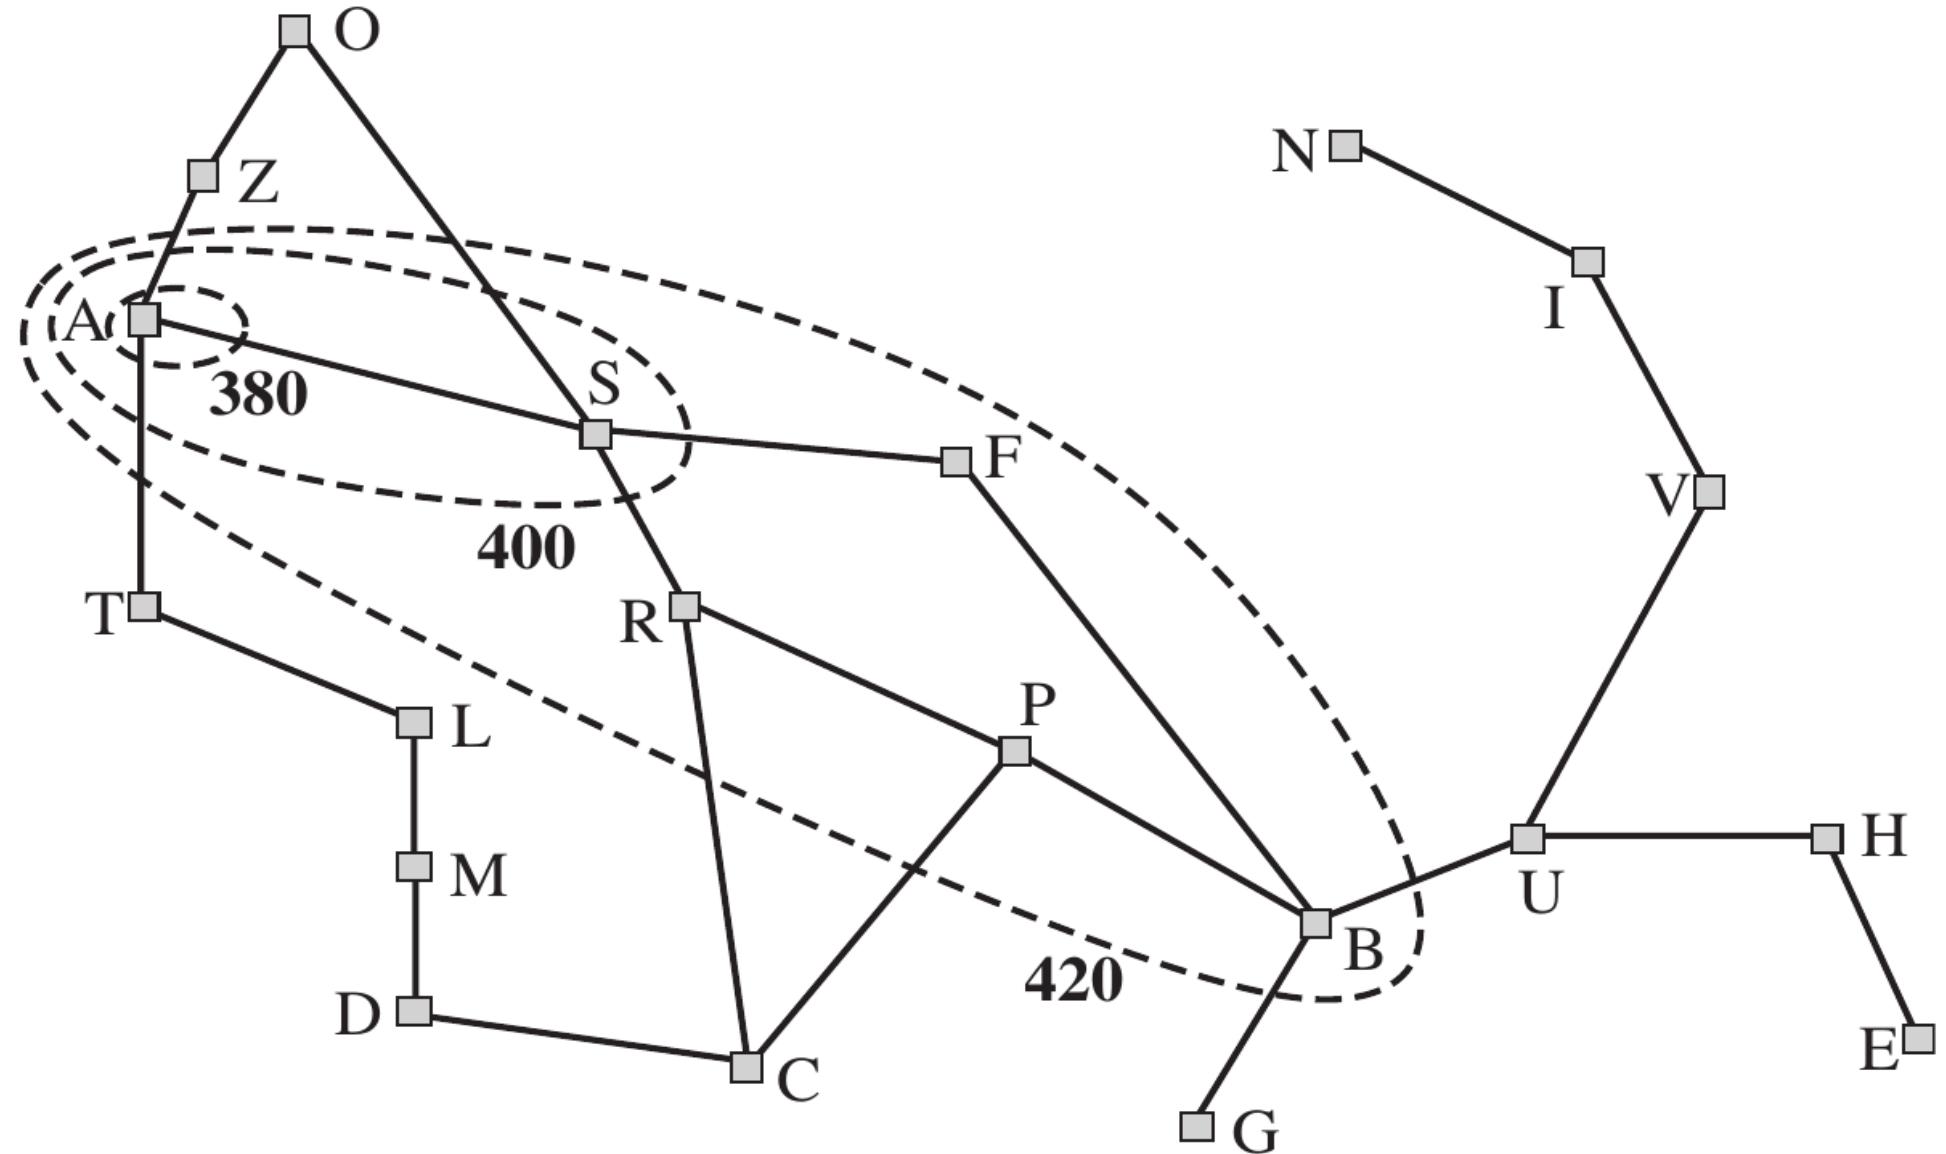
Stages in A* search

Nodes labelled with $f = g+h$

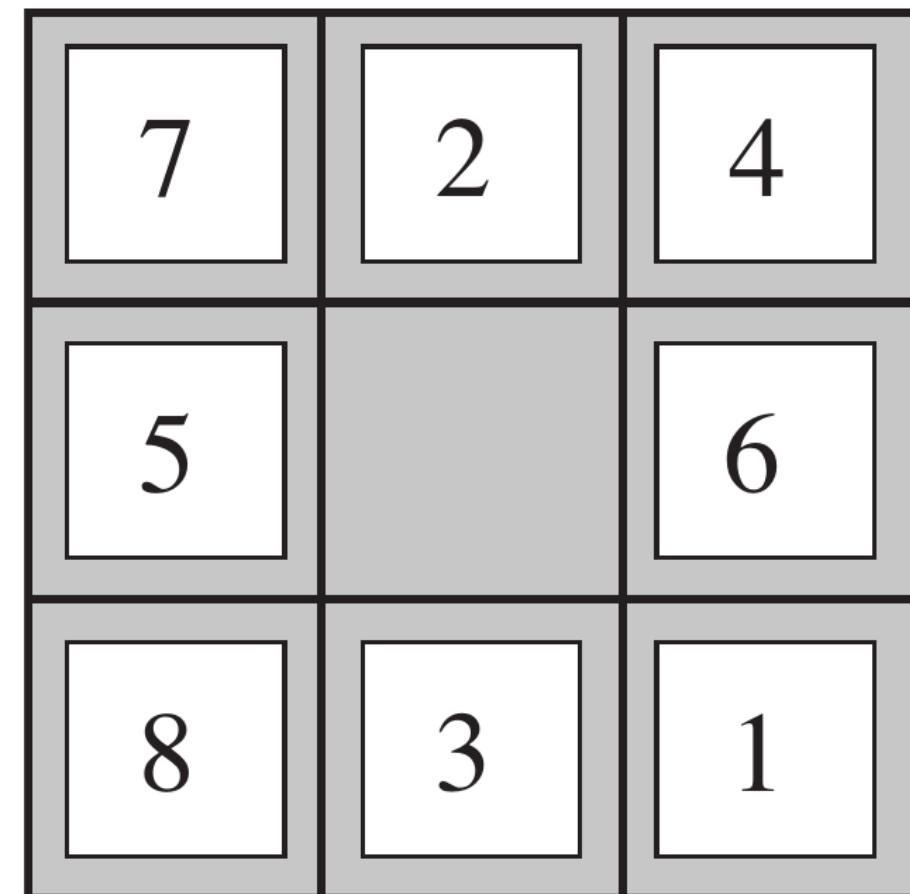
(f) After expanding Pitesti



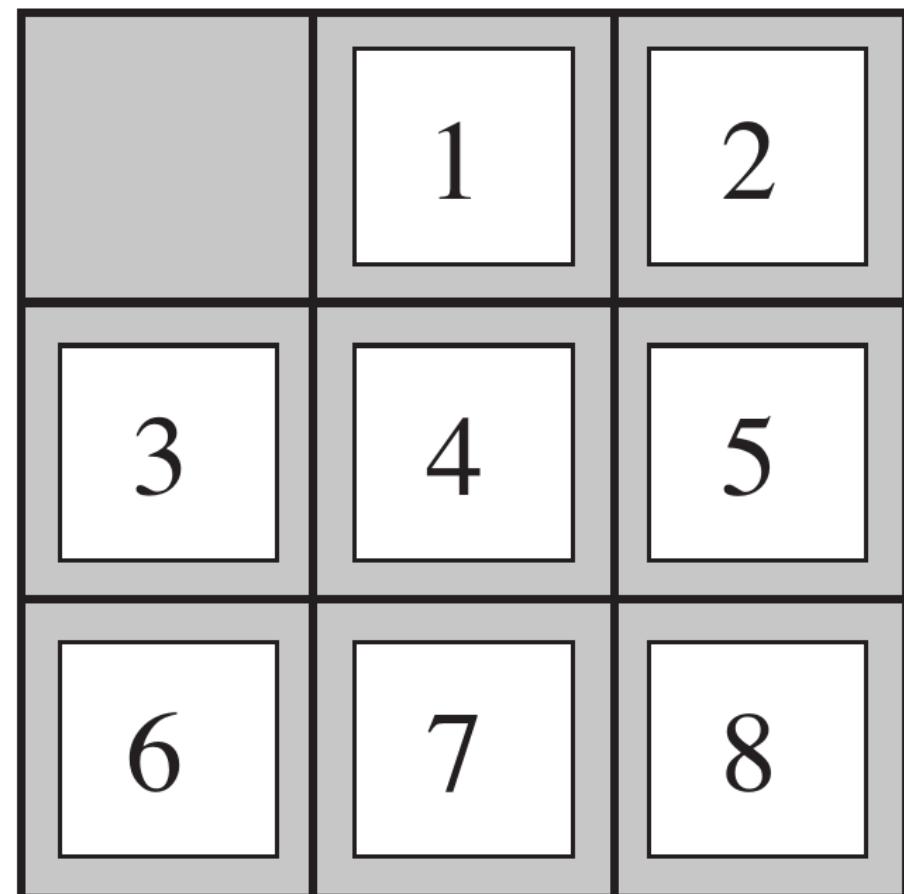
Map of Romania with contours of $f = g+h$



A typical instance of the 8-puzzle



Start State



Goal State

Heuristics for the sliding puzzle

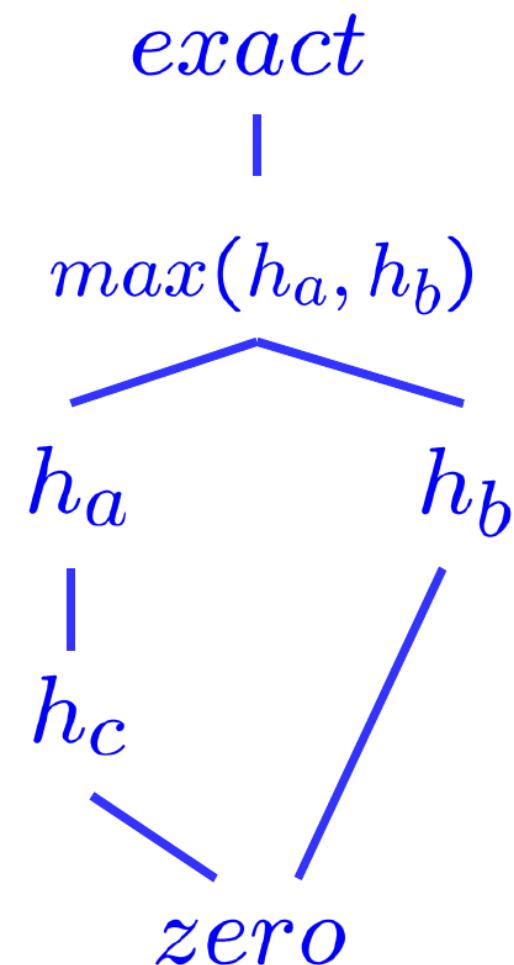
- Heuristic: Number of tiles misplaced
 - Why is it admissible?
 - This is relaxed-problem heuristic
- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?
 - Total Manhattan distance
 - Why is it admissible?

More remarks on heuristics

- How about using the actual cost as a heuristic?
 - Would it be admissible?
 - Would we save on nodes expanded?
 - What's wrong with it?
- With A*: a trade-off between quality of estimate and work per node
 - As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

Trivial Heuristics, Dominance

- Dominance: $h_a \geq h_c$ if
 $\forall n : h_a(n) \geq h_c(n)$
- Heuristics form a semi-lattice:
 - Max of admissible heuristics is admissible
$$h(n) = \max(h_a(n), h_b(n))$$
- Trivial heuristics
 - Bottom of lattice is the zero heuristic (what does this give us?)
 - Top of lattice is the exact heuristic



Optimality of A* Tree Search

- A* tree search is optimal if the heuristic is admissible
- UCS is a special case ($h=0$)